

COMP813 Big AI Project

A Computational Cognitive Model for General Purpose Autonomous Learning Machines

Stone Fang (Student ID: 19049045)

Contents

1	Part A: Critical Review	2
1.1	Review of Recent Work	2
1.1.1	Visual Physics Reasoning	2
1.1.2	Unsupervised Intuitive Physics	3
1.1.3	Make Meanings from Sensory Input	4
1.1.4	Curiosity Driven & Self-aware Agents	4
1.2	Critical Discussion: Cognitive Perspective	5
1.2.1	Evidences from Cognitive Science	5
1.2.2	Visual Physics Reasoning	5
1.2.3	Unsupervised Intuitive Physics	6
1.2.4	Make Meanings from Sensory Input	6
1.2.5	Curiosity Driven & Self-aware Agents	6
2	Part B: A Cognitive Learning Model	7
2.1	Model Description	7
2.1.1	Principles and Hypothesis	7
2.1.2	Algorithm	7
2.1.3	Choice of the Form of Φ	9
2.2	Potential Applications	9
2.2.1	Object Detection and Tracking	9
2.2.2	Object Recognition	9
2.2.3	Image Understanding and Visual Reasoning	10
2.2.4	Knowledge Accumulation	10
2.3	Conclusion and Future Works	10
	Reference	12

1 Part A: Critical Review

This part is a review of recent studies on computational infant learning models and theories. The review is focused on visual physics intelligence such as object tracking and physical reasoning because it is the fundamental of human intelligence (Goswami, 2019). First the main ideas of the studies are briefly described, and then a critical review will be conducted from cognitive science perspective.

1.1 Review of Recent Work

1.1.1 Visual Physics Reasoning

The paper of (Riochet et al., 2018) mainly has two parts. In the first part, the authors proposed IntPhys is a benchmark to diagnose AI systems on visual intuitive physics reasoning tasks, inspired by studies of intuitive physics on infants. In the second part, the authors proposed two unsupervised deep neural networks “infant” learning models on intuitive physics.

1.1.1.1 Diagnostic test benchmark The first part aims at answering a basic problem: evaluation. That is, how can we be sure that a system has certain level of understanding of physics? The authors argue that end-to-end visual tasks such as 3D structure recovery, object tracking, or visual question answering (VQA) are not suitable for such evaluation because of 1) dataset bias and 2) noise measure. For example, a VQA system does not perform well could be because of, not poor understanding of physics, but a bad language model. In order to solve this problem, the IntPhys benchmark is proposed, containing a set of videos as “unit tests” which are independent of any end-to-end task. The tests are organised into four categories: 1) Object permanence, 2) Shape constancy, 3) Spatio-temporal continuity, and 4) Energy/Momentum. It is inspired by the of “violation of expectation” revealed by psychologists, that is, an infant or animal will be “surprised” by visual scenes which is physically impossible. Following this idea, the “physical plausibility” score is introduced to measure the level of surprise. This score is expected to be lower on video clips violating physical laws than that comply.

1.1.1.2 Two “infant” learning models In the second part, two “infant” learning models are proposed. They are both unsupervised/self-supervised neural network models to learn intuitive physics merely from visual presentations in first-person viewpoint. For the sake of simplification, the models are not allowed to interact with the settings, though this is not the real situation of infants or animals. The first model is a Convolution Neural Network (CNN) encoder-decoder model of ResNet-18 pre-trained on ImageNet, and the other is a conditional Generative Adversarial Network (GAN). The models are trained on the objective of next frame prediction, so there is no requirement for data labelling. The two models are tested on both short-term (5 frames) and long-term (35 frames) tasks. Then, the plausibility score for a frame is computed on the basis of comparison between the prediction \hat{f}_t and the ground truth f_t .

1.1.2 Unsupervised Intuitive Physics

An remarkable nature of infant learning is the fact that physical representations and laws are learned without explicit supervision. To mimic this functionality on machines, a recent study of (Ehrhardt, Monszpart, Mitra, & Vedaldi, 2018) proposed an approach to predict meaningful physical parameters, such as object position and velocity, from raw visual representation of real data without any unsupervised or simulator. Instead of mere prediction of future frames, the goal of this study is to construct an intelligent agent that can learn physical states such as positions and velocities of objects, as well as physical laws enabling the prediction of changes of such states (not changes of appearance) over time.

The solution consists of two steps. The first step is to build a tracker that learns to discover objects and extract the positions. The object detection problem is formalised as to learn a function $\Phi(\mathbf{x}_t) = u_t \in \mathbb{R}^2$ where $\mathbf{x}_t \in \mathbb{R}^{H \times W \times 3}$ is an RGB video frame containing a single object and u_t is the extracted 2D position of the object. This function is learned without any label or a-prior knowledge about that object. More concretely, the function $\Phi(\mathbf{x}_t)$ is modeled as follows:

1. calculates a scalar score $f_v \in \mathbb{R}$ at each pixel v with a shallow Convolutional Neural Network (CNN), generating a heat map.
2. f_v is normalised to a probability distribution s_v by softmax function.
3. the location u is extracted as the expectation over s_v , that is, $u = \sum_v v s_v$.

Then, the algorithm to learn Φ consists of the following parts:

1. causality principle: in a video of real physical process, the trajectory of objects are causal and smooth, that is, physically plausible; on the other hand, if the frames in a video are randomly shuffled, the trajectory should be not. Therefore, a discriminator network $D(\Phi)$ is incorporated that can classify between true physical ordering of frames and random shuffled ones, which results in the classification loss of \mathcal{L}_{disc} .
2. equivariance principle: this principle claims that if a frame \mathbf{x}_t is applied by a transformation g , then the detector's output should be equivalent to be applied by the same transformation, that is, $\Phi(g\mathbf{x}_t) = g\Phi(\mathbf{x}_t)$. This is called a Siamese branch and the training loss is noted as \mathcal{L}_{siam} .
3. penalty on distribution: Because s_t is modeling the distribution of the object's position, it is reasonable to expect it have a peaky shape. To encourage the sharpness of the distribution, the entropy is taken as the penalty adding to the overall loss: $\mathcal{L}_{ent} = -\sum_{v \in \Omega} s_v \log(s_v)$.

Finally the overall loss is a weighted sum of the three items above: $\mathcal{L} = \lambda_d \mathcal{L}_{disc} + \lambda_e \mathcal{L}_{ent} + \lambda_s \mathcal{L}_{siam}$.

For multiple objects, the detector works in a one-by-one way: first, it detects a single object, and then masks it from the frame, and then repeats the process. However, the proposal assumes the number of objects is given.

This tracker is claimed to be well scalable to large datasets, and robust on a variety of objects in different kinds and shapes without prior knowledge..

Then, the second step is to build a predictor that can predict the positions through time by extrapolation of the positions extracted in the first step.

This work also proposed a dataset of Roll4Real containing videos of balls rolling on different surfaces.

The research team also conducted a research to do similar unsupervised learning from past experience (Ehrhardt, Monszpart, Mitra, & Vedaldi, 2019).

1.1.3 Make Meanings from Sensory Input

A significant characteristic of human learning is emerging symbolic represented concept from raw sensory data. In (Evans, Hernandez-Orallo, Welbl, Kohli, & Sergot, 2019) the authors give an answer to this fundamental question concerning human and machine learning: what is the meaning of “make sense” of raw sensory input. In their study, such intuitive approach is formalised by unsupervised program synthesis.

The authors believe that “making sense” of sensory data means interpreting the data by conceptualised and symbolic representation such as objects, which is more than the prediction of future values. The authors also criticise the neural network to be perceptive but not **apperceive** because, while it is trained well to classify images, it can not integrate such classification with other knowledge and do further reasoning. It is also proposed that, inspired by Kant’s theory of the “synthetic unity of apperception”, a theory that can make sense of data must satisfy four unity conditions, namely spatial, conceptual, static, and temporal.

It also proposed a model called “Apperception Engine”, which satisfies such constraints and is able to learn from small amount of data thanks to its strong inductive bias. This is achieved by introducing a casual programming language, Datalog, and synthesizing a Datalog program from sensory input that explains the input and also complies with the four unity conditions. Therefore, the inductive bias comes from two aspects: specification of Datalog language and Kantian unity conditions. Datalog is based on Inductive Logic Programming and the Kantian conditions are expressed by formal logical expressions.

This model is designed as a *general purpose* apperception system, and has been tested on a range of domains. It is claimed to significantly out-perform neural network baselines, and particularly, to reach human-level in the sequence induction IQ tests, which would be more notable considering the system is not specifically designed to solve such tasks.

1.1.4 Curiosity Driven & Self-aware Agents

Infants are experts at generating orderly action from unstructured settings by playing without any supervision or even explicit external rewards. Inspired by this, a study (Haber, Mrowca, Wang, Fei-Fei, & Yamins, 2018) was conducted to model such ability with an agent which is intrinsically motivated by “implementation” of curiosity and self-awareness. In this study, an agent are designed to learn how to play like an infant via deep reinforcement learning approach.

The agent is constructed by two interacting models: a world-model and a self-model. The world-model is a predictor of physical events based on its observational inputs, while the self-model is an estimator of the loss of world-model based on both the input and potential actions. The action of the agent is always chosen to maximise the loss of the world-model, that is, antagonize the world-model. To aim this, the self-model is introduced to estimate such loss, which is claimed to be a representation of self-awareness on its internal state, and a mathematical formalisation

of “interestingness” and “curiosity”. The two components are both modeled by convolutional neural networks.

In the execution of this algorithms, the agent always pays attention to more challenging objects once it has learned some knowledge. In other words, it gets “bored” on the things that it has already “understood” well and always tries to challenge itself, resulting in a curiosity-driven behavior.

In the future work, the authors are planning to build computational models more than a robust learning agent, but precisely quantitatively comparable to human children’s development.

1.2 Critical Discussion: Cognitive Perspective

This part criticises the learning algorithms mentioned above from a cognitive perspective. In this paper, an AI algorithm is not evaluated by how well it performed on certain tasks as most AI researchers do; instead, it is criticised from the cognitive science viewpoint. To be precise, it will be evaluated on whether it provides insights into the mechanism of cognitive process. As demonstrated in the work of (Yeap, 2011), the mapping algorithm is examined by comparing to three characteristics of cognitive mapping: fragmented, incomplete and imprecise. Similar idea is also expressed in (Dupoux, 2016), in which a reverse engineering approach is proposed in order to, from AI studies, gain scientific insights about underlying mechanisms of humans’ psychological process, especially language acquisition.

1.2.1 Evidences from Cognitive Science

Studies in cognitive science uncovered some characteristics of humans’ learning. Comparing to AI, the natural intelligence differs in the following aspects significantly:

- structured representation: it is proven that infants have the concept of objects, understand the physical law such as object permanence, and can organise raw information into categories. (Goswami, 2019; Stangor & Walinga, 2014)
- absence of supervision and reward: infants learn the physical concepts and rules without explicit supervision and reward; instead, natural intelligences build internal representation of physics inductively and automatically from raw perceptual data and experience. (Ehrhardt et al., 2018; Stangor & Walinga, 2014)

1.2.2 Visual Physics Reasoning

In work of (Riochet et al., 2018) as reviewed in section 1.1.1, two unsupervised “infant” learning models are proposed. The advantage of these models is the absence of supervision or external reward, which is more likely as the situation of real infant learning. However, these proposals are still task-oriented models, because they are trained to learn frame predication on raw pixel inputs. As a result of the “black-box” natural of neural network, it is hard to tell what is inside such model and, more essentially, whether the model learns the concept of “object” and the law of “object permanence” as humans do. Therefore, we can gain little inspiration from this “black-box” system on the cognitive process of how human infants learns. On some extent, learning predictions on visual sequences is a similar task to learning a language model on raw texts as

word sequences, which still have difficulties to capture the complex structures and reasoning in humans' language.

1.2.3 Unsupervised Intuitive Physics

The study reviewed in section 1.1.2 is a novel approach in intuitive physics learning. It assumes the form of learning function Φ , which is a form of knowledge on the concept of object. That is, for this algorithm, the concept of object is modeled as a-prior knowledge instead of learnable variable from data.

Furthermore, the proposal assumes the number of objects is given on multiple objects positioning, which is a drawback because this is a specific instead of general assumption, and makes the algorithm less practical in real world applications. From cognitive perspective, human learns the concept of object and the ability of detection first, and the the number of objects second. Therefore, when a infant learns the detection of object, it is impossible to know the number of objects in advance.

1.2.4 Make Meanings from Sensory Input

The model reviewed in section 1.1.3 is unsupervised and independent of domain knowledge, which are significant advantages. The core of it is the Datalog language which provides strong inductive bias. Datalog is not embedded by any prior concepts; instead, all it preset rules is its logical form and four unity constraints, and the concepts are learned from input data. It was tested in a variety of domains and performed well, which in some extent demonstrates its nature of "general intelligence".

However, the Datalog language is a strong bias, it equivalently presets strong rules on the model, which have not been proven by cognitive studies. Furthermore, though the model was tested on different domains, they are all tasks more relying on logic and reasoning which is suitable for a logic based system. By contrast, humans can do both logical and intuitive tasks with one single system – the brain. The model may need more tests on intuitive tasks such as face recognition and image segmentation.

1.2.5 Curiosity Driven & Self-aware Agents

The study reviewed in section 1.1.4 is different from the others above because it focuses on the "motivation" aspect of a learning system. In a addition to a "world-model" modeling the outside world, the AI agent employs a "self-model" to model the inner state as a simulation of humans' curiosity and self-awareness. With this approach, the agent can learn without any extrinsic supervision or reward.

From the perspective of cognitive science, however, the biggest concern on this approach is the machine equivalent of humans' curiosity and self-awareness. There has not been sufficient evidence to prove that the algorithm works the same way as humans' cognitive process. The "self-model" estimates the loss of "world-model" and thus provides the knowledge of inner state, but curiosity and self-awareness are much more than estimation. To prove the self-awareness, at least a mirror test (Wikipedia contributors, 2019c) should be performed on this agent.

2 Part B: A Cognitive Learning Model

In this part, an artificial will be discussed aiming at inspiring the understanding on human infants' learning process. First, a cognitive intelligent algorithm for object segmentation and tracking will be introduced and analysed following the approach in (Yeap, 2011). Second, a variety of questions will be investigated by this approach.

2.1 Model Description

As summarised before, human infants learning in early stage is entirely unsupervised, without explicit supervision or extrinsic reward. As a result, the proposed approach should also have these characteristics. To minimise any prior knowledge implied in the model, the input view of an infant is the only starting point, and the bias is preferred to more general hypothesis rather than specific ones.

2.1.1 Principles and Hypothesis

The following principles form the base of the algorithm:

- **Association:** or correlation/relativeness. It is a fundamental cognitive process initially identified by Pavlov's famous experiment of dogs on classical conditioning (Stangor & Walinga, 2014). Though it is mostly used to describe structured behaviour in psychology, in this paper this principle is extended to a universal extent. That is, if some things often occur simultaneously, they might also do so in the future. The more frequent their occurrence happens in the past, the more confident we assert it in the future.
- **Attention:** or selectiveness. Attention is a necessary condition of visual recognition (Goswami, 2019). With this cognitive mechanism, human can concentrate on a selected discrete aspect of information while ignore others (Wikipedia contributors, 2019b). This is essential because it prevents the human brain from disorientation in the unmeasurable amount of sensory input.
- **Abstraction:** or conceptualisation. Abstraction is one of the key characteristics of human behaviour (Wikipedia contributors, 2019a). It is the way of hiding information and thus managing complexity (Abelson & Sussman, 1996). Abstraction can also be seen as the process of compression, which is an implementation of minimum description length (MDL) principle (Henderson & Muggleton, 2014).

The principles introduced above is quite fundamental and general in human learning behaviour. As a result, the model following these principles is expected to have minimal inductive bias.

2.1.2 Algorithm

The input of the model is time sequences of raw pixel images $\mathbf{x} = \{x_t\}$. For the sake of simplicity, it is assumed that there is only one object moving in the scene. This is reasonable because infants always start their learn from simple views.

1. Attention

Extract an area which changes across frames by simply calculate the difference between two consecutive frames $\Delta x_t = x_t - x_{t-1}$. Obviously, in common cases, the changing area is caused by moving objects. However, no prior concept of “object” or “motion” is introduced into this model; instead, only “change” is a presumption and the concepts of “object” and “motion” are learned. The area having changed between frames is the attention area.

After the changing area has been detected, match this area between the two frames by offset and comparison, resulting in the detection of the matched boundary m_t in each frame.

2. Association

After the matched boundary m_t has been extracted, the model will build an representation or concept by a encoding function Φ : $c_m = \Phi_m(m_t)$. According to the association rule, the encoding function Φ should also give a confidence score, depending on how frequently the agent has observed the same view in the matched boundary m_t in the past.

3. Abstraction

The next step is re- calculating and structuring the representations or concepts to emerge new abstractions. In this step the agent should be able to:

- Merge similar concepts into categories to emerge abstract concepts. This is to form generalised concept from individual samples.
- Decompose a concept into small pieces to shorten the encoding. Because the attention area is merely calculated by the difference between frames, it may contains several objects or a complex object decomposable into small ones. The agent cannot be aware of this fact; instead, it builds concepts autonomously from the view. Therefore, in order to merge the right concepts of objects, the model should be able to decompose a representation into small parts, each of which could be a single object or a sub-component. However, the “small part” is not created by arbitrary decomposition; instead, it should follow the association principle. That is, it must reach a minimum frequency of occurrence as the support to become a concept.
- Compose several concepts to form compound concepts. With this approach the agent can emerge more complex and abstract concepts. Such composition also follows association principle: the more frequently some concepts are observed concurrently, the more likely they are associated and form a compound concept.

At this step the model will form a set of concepts $\mathcal{C} = \{c_i\}$ with corresponding encoding function $\Phi = \{\Phi_{c_i}\}$ where $c_i = \Phi_{c_i}(\mathbf{c}, \mathbf{x})$, $\mathbf{c} \subset \mathcal{C}$, which form the knowledge of the agent $\mathcal{K} = \{\mathcal{C}, \Phi\}$

These three steps will be executed iteratively. After the agent has already accumulated some knowledge, it will apply the encoding functions onto a new observation, resulting in detected objects which it has already learned. Then more “attention” will be paid to unknown area while the known areas are more likely to be ignored.

2.1.3 Choice of the Form of Φ

In section 2.1.2 the algorithm is introduced without a specific form of concept encoding function Φ . Rather than that, only some constraints are set. To make the model more convincing as well as more practical, the possible forms of encoding function Φ are discussed in this section.

The Bayesian Program Learning(BPL)(Lake, Salakhutdinov, & Tenenbaum, 2015) is the best solution that I can find so far to the encoding function. Following this method, concepts are encoded by stochastic programs, which are defined as probabilistic generative models described as statements of structured programming languages. The reason of its suitability is twofold. On one hand, it is a probabilistic model, which means it aligns essentially with the associative nature of the model; on the other hand, it is represented in the form of structured language, which means it can be composed and decomposed easily thanks to its symbolic and recursive nature, with an additional benefit of explainability.

However, the stochastic programs in the generative model are in fixed structure instead of learned. This may be improved by the abstraction invention algorithm proposed in (Henderson & Muggleton, 2014). In this framework, new abstractions are created from existing “background knowledge” by inverse β -reduction. This method is within a larger inductive programming scope of Compression-Based Learning (CBL), which learns by finding the minimum length to describe the input data.

2.2 Potential Applications

The proposed model can be used in a range of intelligent tasks. Because of the nature of both logic and statistics, the model is suitable for both intuitive and reasoning tasks.

2.2.1 Object Detection and Tracking

The model can be used for object detection tasks in either image or video. By design, the model pays more attention to moving objects so it is also suitable for dynamic object tracking. The model learns a set of concepts and corresponding encoding functions, which can be directly applied to a image or video frame to detect objects of which it has the knowledge. However, this approach results in a exhaustive search in a high computational complexity of $O(SN_c)$ where S is the size of image and N_c is the number of concepts, which is not likely the case of what humans do more smartly. To cope with this, it is optimised by structuring all the concepts into a hierarchy, and then only running the lowest level encoding functions. If a primitive concept was not detected, then it is not necessary to run any of the compound concept containing that primitive concept.

2.2.2 Object Recognition

The model can be applied in object recognition tasks. Note the model continuously executes concept refinement and abstraction, so it will consequently improve its representation on the classification of objects and become more tolerant on noise and variation. In addition, a significant characteristic of this model is that the capabilities on object recognition and detection are not separable – to recognise an object one has to detect it first; however, to some extent, this

model detects an object by recognising it. This seems paradoxical at first sight but if you think it twice it is exactly the same as humans' cognition – humans can detect an object more easily if they are more familiar with it; on the contrary, it is more difficult to detect an object that has never been seen. A typical example of this phenomenon is to read the characters in some language. A human can detect the symbols in a language he or she speaks effortlessly, but not so easily to detect a symbol from an unfamiliar language. Though the human can also do the detection by some basic rules such as different symbols are approximately in the equal size, or the lines and curves in one symbol are more cohesive than that in different symbols, but he or she needs to pay more attention to that symbol and spend more mental resources, merely resulting in a less confident output.

2.2.3 Image Understanding and Visual Reasoning

Thanks to its structured and logical nature, this model is suitable for visual understanding and reasoning tasks, for example, image captioning, physical state prediction, and so on. Through the learning process the agent can gain fundamental ability for visual intelligent tasks. To achieve better performance, the model needs to be fine-tuned on specific tasks. However, this model has two advantages on visual reasoning tasks. One is that some basic concepts that the model learned are transferable because of its fundamentality and generality, thus it needs less training samples to learn any physical law. Another benefit is its symbolic and logical representation, which would be more easily to incorporate existing knowledge. This advantage can also significantly reduce the size of training data it requires, leading to one-shot or even zero-shot learning.

2.2.4 Knowledge Accumulation

This model can accumulate knowledge continuously with long-term learning. Even without an explicit memory, this model memorises the information it has ever seen and learned in the posterior probabilities of the bayesian generative models. This may also be improved by giving a larger weight to more recent observations so that the model can mimic humans' forgetting process. If the model had a powerful learning ability, via continuous learning of knowledge from the observations of the nature, it could grasp basic common sense regarding the physical world underpinning humans' knowledge and ultimately reach humans' level of intelligence, in the meaning of cognitive processes.

2.3 Conclusion and Future Works

To summarise this part, a cognitive learning model was proposed for visual intelligence. This model follows three principles of attention, association, and abstraction, resulting in autonomous learning without explicit supervision, external reward, or any a-priori knowledge. This model is suitable for a variety of intelligent tasks including visual object detection, tracking, recognition, reasoning, as well as accumulating knowledge in long-term learning.

However, there is still some problems to be solved. First, the attention mechanism needs to be refined. In this article, attention is paid to two kinds of areas: one containing changes and the other containing unknown objects. However, it needs more investigation whether more attention method is needed. Furthermore, if different attention methods can be unified, the model would

be more simple and mathematically beautiful. Second, in terms of the association principle, it claims that the observations with frequent concurrence form a concept, but does not give a quantitative solution of the frequency threshold for creating new concepts. The threshold is a hyperparameter in this model, which need to be justified on the choice of its value. Furthermore, it would be a better solution if this hyperparameter can be eliminated. Third, the abstraction representation framework and learning algorithm need to be formalised in mathematics for rigorousness and executability, and more effective implementations need further investigation for scalability on large dataset.

Reference

- Abelson, H., & Sussman, G. J. (1996). *Structure and interpretation of computer programs* (2nd ed.). Cambridge, MA, USA: MIT Press.
- Dupoux, E. (2016). Cognitive science in the era of artificial intelligence: A roadmap for reverse-engineering the infant language-learner. *CoRR*, *abs/1607.08723*. Retrieved from <http://arxiv.org/abs/1607.08723>
- Ehrhardt, S., Monszpart, A., Mitra, N. J., & Vedaldi, A. (2018). Unsupervised intuitive physics from visual observations. *CoRR*, *abs/1805.05086*. Retrieved from <http://arxiv.org/abs/1805.05086>
- Ehrhardt, S., Monszpart, A., Mitra, N. J., & Vedaldi, A. (2019). Unsupervised intuitive physics from past experiences. *CoRR*, *abs/1905.10793*. Retrieved from <http://arxiv.org/abs/1905.10793>
- Evans, R., Hernandez-Orallo, J., Welbl, J., Kohli, P., & Sergot, M. (2019). *Making sense of sensory input*. Retrieved from <http://arxiv.org/abs/1910.02227>
- Goswami, U. (2019). *Cognitive development and cognitive neuroscience : The learning brain*. Routledge. Retrieved from <http://ezproxy.aut.ac.nz/login?url=http://search.ebscohost.com/login.aspx?direct=true&db=cat05020a&AN=aut.b27297470&site=eds-live>
- Haber, N., Mrowca, D., Wang, S., Fei-Fei, L., & Yamins, D. L. K. (2018). Learning to play with intrinsically-motivated, self-aware agents. *Proceedings of the 32Nd international conference on neural information processing systems*, 8398–8409. Retrieved from <http://dl.acm.org/citation.cfm?id=3327757.3327931>
- Henderson, R., & Muggleton, S. (2014). Automatic invention of functional abstractions. In *Latest advances in inductive logic programming* (pp. 217–224). https://doi.org/10.1142/9781783265091_0023
- Lake, B. M., Salakhutdinov, R., & Tenenbaum, J. B. (2015). Human-level concept learning through probabilistic program induction. *Science*, *350*(6266), 1332–1338. <https://doi.org/10.1126/science.aab3050>
- Riochet, R., Castro, M. Y., Bernard, M., Lerer, A., Fergus, R., Izard, V., & Dupoux, E. (2018). IntPhys: A framework and benchmark for visual intuitive physics reasoning. *CoRR*, *abs/1803.07616*. Retrieved from <http://arxiv.org/abs/1803.07616>
- Stangor, C., & Walinga, J. (2014). *Introduction to psychology – 1st canadian edition*. Victoria, B.C., Canada: BCcampus. Retrieved from <https://opentextbc.ca/introductiontopsychology>
- Wikipedia contributors. (2019a). *Abstraction — Wikipedia, the free encyclopedia*. Retrieved from <https://en.wikipedia.org/w/index.php?title=Abstraction&oldid=922555014>
- Wikipedia contributors. (2019b). *Attention — Wikipedia, the free encyclopedia*. Retrieved from <https://en.wikipedia.org/w/index.php?title=Attention&oldid=921297171>
- Wikipedia contributors. (2019c). *Mirror test — Wikipedia, the free encyclopedia*. Retrieved from https://en.wikipedia.org/w/index.php?title=Mirror_test&oldid=921100578
- Yeap, W. (2011). *A computational theory of human perceptual mapping*. Cognitive Science Society. Retrieved from <http://ezproxy.aut.ac.nz/login?url=http://search.ebscohost.com/login.aspx?direct=true&db=ir00946a&AN=auck.10292.3513&site=eds-live>