

Popular Topic Mining from Blogs

Stone Fang (Student ID: 19049045)

Computers and Information Sciences

Auckland University of Technology

Auckland, New Zealand

fkn7060@autuni.ac.nz

I. OVERVIEW

People's concerns and opinions are important reference for innovations of new products or services. However, accomplishing such task by humans is expensive, time-consuming and difficult to scale. As a response, a number of individuals and organisations are leveraging text mining technologies to mining meaningful information from large volume of text such as news media [1]. Among a variety of studies and applications, topic modelling is an important method to extract hot topics which reflects public attention and opinion from massive texts [1]–[3]. However, effective method of extracting useful information from text on the Internet remains an open challenge [3].

Evaluation of topics mined from text is another challenge, mostly due to the lack of ground truth because topic modelling is an unsupervised learning task [4].

The goal of this project is to mine most popular topics that people were discussing from blog posts by utilising various text mining algorithms and tools. Specifically, we will find two most popular topics for each group in the following demographics:

- Males
- Females
- People 20 years old or younger
- People older than 20
- Everyone

The remainder of this article is organised as follows. In ?? related works on topic mining and evaluation will be reviewed. The methodology of topic mining are detailed in section III, while the results, analysis and evaluations are presented in section IV. The works of this article are summarised in section V and open issues and future works are discussed in section VI.

II. LITERATURE REVIEW

Jacobi, Atteveldt, and Welbers [1] conducted an in-depth study of how to apply topic modelling technologies on analysis of qualitative data in academic research.

Boyd-Graber, Mimno, and Newman [4] provides a summary of topic evaluation methods, which are divided into three categories: human evaluation, diagnostic metrics, and coherent metrics. The first one needs human effort so it is expensive and time-consuming, while the other two can be calculated by computer without human interference.

Human Evaluation requires human involvement in the evaluation task. One method in this category is accomplished by word intrusion task. Specifically, a person will be presented by a list of words and is asked to find an intruder in the meaning of not belonging to others. The words list are constructed by first selecting highly possible words from a topic, and then randomly choose one word with low probability in the same topic but high probability in a different topic. If the intruders are easily to be identified, then the topic is more likely coherent [4].

Diagnostic Metrics only compute statistics of topics without requirements of external knowledge source. Some methods in this category are [4]:

- **Topic Size:** measured by the sum of numbers of tokens belonging to a certain topic. Generally speaking, small topic size means low quality.
- **Word Length:** average length of N most dominant words in a topic. The usefulness of this metric is corpus dependent.
- **Corpus Distribution Distance:** A probability distribution can be derived from a topic over the vocabulary, and further normalised by global word count in the whole dataset. The distance between different topics reflects how much these topics are separated.

Coherence Metrics is a type of methods which automatically compute score of topic coherence, and their accuracy is close to human performance. The basic idea is measuring how a pair of words from top N dominant words are associated [4]. It is formalised as

$$TC-f(\mathbf{w}) = \sum_{i < j} f(w_i, w_j), i, j \in \{1 \dots N\}$$

where $\mathbf{w} = \{w_1, w_2, \dots, w_N\}$ is the list of N most dominant words, and f is the scoring function of association between two words. There are a variety of ways to compute f , such as counting the co-occurrence of two words, or counting the number of documents containing both words. Two popular implementations of f are pointwise mutual information (PMI) and log conditional probability (LCP) [4].

$$PMI(w_i, w_j) = \log \frac{P(w_i, w_j)}{P(w_i)P(w_j)}$$

$$LCP(w_i, w_j) = \log \frac{P(w_i, w_j)}{P(w_j)}$$

III. RESEARCH DESIGN

In this section the solution will be described in detail. First an overview of the dataset is given, and then the algorithm of topic mining is detailed.

A. Data Description

The dataset contains 19,320 files in XML format, each containing articles of one person posted generally between 2001 and 2004. Metadata of the bloggers includes gender, age, category, and zodiac. In addition, the number of posts for each person are also counted. The result is summarised by Fig. 1, which is created by Python packages `Pandas` and `Matplotlib`. From this figure we can acquire some basic statistics of the dataset, including

- 1) Gender: data samples are quite evenly distributed over both genders.
- 2) Age: most bloggers are younger than 30, almost of them under 20. On the other hand, there are two gaps around 20 and 30 which may implies some missing data points in the dataset
- 3) Zodiac: The distribution over zodiac is reasonable even.
- 4) Category: the most frequent category is unknown, which is trivial, while the second frequent one is student, far more than other categories.
- 5) Number of posts: most bloggers published less than 100 posts, while the peak appears at 10, which implies people are most likely to write around 10 posts.

B. Topic Mining Algorithm

The general idea for mining popular topics used in this project is to find the most significant “things” mentioned in the overall dataset, as well as the closely related information.

The overall architecture of the algorithm is shown as Fig. 2, and the details of each step are described in the following subsections.

1) *Data Cleaning*: Before applying any text mining techniques, it is important to do basic data cleaning to improve data quality. In this step, a few operations for preprocessing will be carried out based on the observations of the dataset, with details as follows.

- **Problem**: At some place there is no whitespace between a punctuation and the word following it, which causes wrong tokenisation. Specifically, the punctuation might be tokenised with the following word as one token.
Solution: Add whitespace after a punctuation if a word immediately follows it.
- **Problem**: Two or more consecutive quote symbols may cause wrong tokenisation.
Solution: Replace two or more quotes as a double quote.
- **Problem**: The unicode quote may affect tokenisation and stopwords matching.
Solution: Replace unicode quote by ASCII quote.
- **Problem**: The unicode quote may affect tokenisation and stopwords matching.
Solution: Replace unicode quote by ASCII quote.

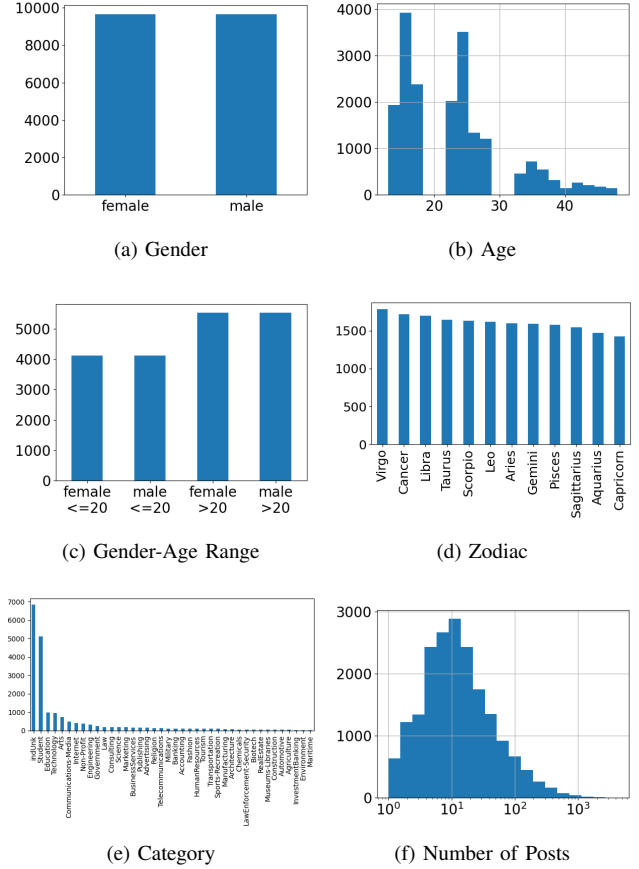


Fig. 1. Data Overview. Histogram over (a) gender (b) age (c) gender and age range (d) zodiac (e) category (f) number of posts

- **Problem**: Characters that are usually not part of normal English text may disturb tokenisation and POS tagging.
Solution: Remove invalid characters such as “*”, “#”, and so on.
- **Problem**: Sometimes people repeat a certain letter in a word for emphasis, but it will result in wrong words and also increase the vocabulary size.
Solution: No English word has more than two consecutive appearances of the same letter, so three or more repetition of a letter is squeezed into two.
- **Problem**:
Solution:

These operations are implemented by regex matching and substitution, or simple text replacing. To use regex, the Python’s `re` package are imported.

2) *Tokenisation*: Tokenisation is usually the first step of all text mining pipelines, which includes sentence and word tokenisation. Sentence tokenisation is to split the whole text into sentences, while word tokenisation splits a sentence into word or tokens. In this project we use `nltk` package to do such task. This package provides two functions `sent_tokenize()` and `word_tokenize()` for both tokenisation. A document is first tokenised into sentences, and then each sentence is

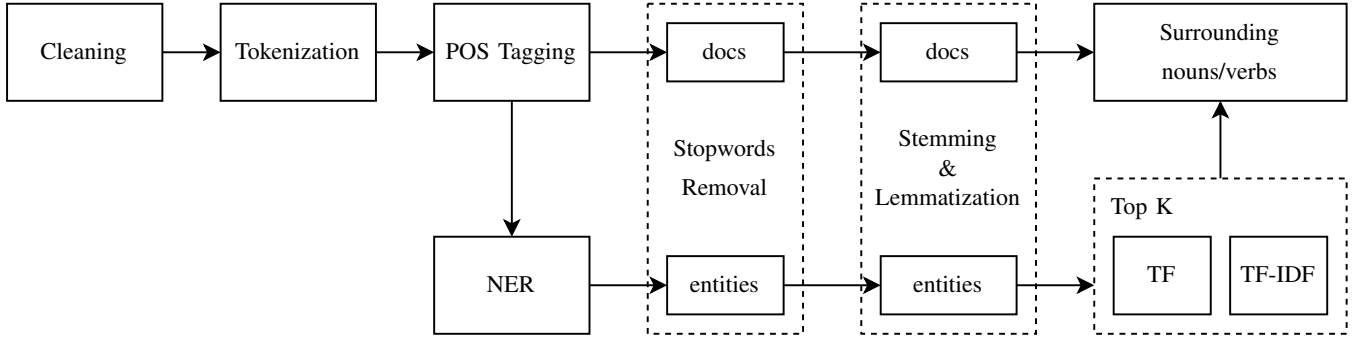


Fig. 2. Overall Architecture of Topic Mining Algorithm

tokenised into words. Finally, a document is represented as list of lists, as each sentence is a list of words.

3) *POS Tagging*: Part-Of-Speech (POS) tagging is the second step following tokenisation. In this step, each word is assigned by a POS tag. `nltk` provides a handy function `pos_tag()` to do this task. This function works on sentence level, and maps each word into a tuple which is the pair of word and POS tag.

4) *Entity Extraction*: In this project, a topic is defined as a “thing” or “object”. Therefore, in order to find the topics, we need to find all “things” or “objects” first. There are a few options to do this task, among which two methods will be employed by this project: Named Entity Recognition (NER) and parsing.

a) *NER*: Named entities are ideal candidates of topics as they denote real-world objects. `nltk` provides a function `ne_chunk()` to extract entities from sentences. The input of the function should be a list of tokens with POS tags, which is another reason why POS tagging should be done in previous step. The return value of this function is a list of chunks, each of which is basically a list and may contain a `label` attribute if it is a recognised entity. The entity type can be acquired by `label()` and the entity itself should be acquired by joining all the elements of the chunk.

b) *Parsing*: Another way to extract objects is parsing by pre-defined patterns. For example, it is reasonable to treat definite nouns as objects according to the grammar.

5) *Stopwords Removal*: Stopwords are most common words which carries no significant meanings. Removing stopwords can reduce the size of data to be proceeded as well as increase the result accuracy. `nltk` provides an out-of-the-box stopwords collection, but the experiment shows that some common words carrying no meaning are not included in the list. In order to expand the stopwords list, more words are collected from website ¹.

Stopwords removal is conducted after POS tagging and entity extraction because these two steps are sequence model, which means their performance rely on word order. If stopwords are removed before them, we will get sentences which do not comply with English grammar. In addition, stopwords

removal are carried out on tagged documents as well as entities extracted. Theoretically, stopwords cannot be entity, but errors will happen in any POS tagging and NER model. Therefore, trying to remove stopwords can reduce the error introduced in previous steps.

6) *Stemming and Lemmatisation*: Stemming and lemmatisation are both techniques for text normalisation, that is, convert an inflected word into its root form. However, stemming and lemmatisation work in different way. Stemming removes suffix or prefix from a word, returning a word stem which is not necessarily a word. On the other hand, lemmatisation always looks for the lemma from word variations with morphological analysis. For example, stemming against the third-person singular form “flies” returns “fli”, while lemmatisation returns “fly”. In this project, these two methods are combined together to reach the maximum extent of word normalisation.

`nltk` provides various stemming algorithms such as `PorterStemmer` and `LancasterStemmer`, and one lemmatisation algorithm `WordNetLemmatizer`. In the code we use `WordNetLemmatizer` followed by `PorterStemmer`.

7) *Word Count and TF-IDF*: After all “objects” have been extracted and normalised, the next step is to find most popular ones as the most dominant topics. Popularity can be defined in various ways, and in this project two approaches are used: word count and Term Frequency-Inverse Document Frequency (TF-IDF). In the first method, we simply count the appearances of each entity and get the most two frequent ones. In the second method, we calculate the TF-IDF value of each entity word, following the definition

$$\begin{aligned} \text{TF-IDF}(t_i, d_j) &= \text{TF}(t_i, d_j) \times \text{IDF}(t_i) \\ &= \text{TF}(t_i, d_j) \log \frac{N}{\text{DF}(t_i)} \end{aligned}$$

$\text{TF}(t_i, d_j)$ is the Term Frequency of term t_i in document d_j , which is computed by count of t_i in d_j divided by the total number of terms in d_j . $\text{DF}(t_i)$ is the Document Frequency, which is the number of documents that contains t_i . As we can see here, TF-IDF is a term-document-wise number so a term has different TF-IDF values in different documents. In order to rank all terms over the whole dataset, TF-IDF values of

¹<https://gist.github.com/sebleier/554280>

Two different methods might return different results, which will be compared and analysed in section IV.

C. Evaluation Method

IV. RESULT, ANALYSIS, AND EVALUATION

A. Result

B. Analysis

V. CONCLUSION





VI. OPEN ISSUES AND FUTURE WORKS

fake wtdwt dwt wd twdt w d dw w wdt wd tw dtdw wd
dwt d tagfdygafdgkdwlrn gdwgd fake wtdwt dwt wd twdt w
d dw w wdt wd tw dtdw wd dwt d tagfdygafdgkdwlrn fake
wtdwt dwt wd twdt w d dw w wdt wd tw dtdw wd dwt d
tagfdygafdgkdwlrn

fake wtdwt dwt wd twdt w d dw w wdt wd tw dtdw wd
dwt d tagfdygafdgkdwlrn

fake wtdwt dwt wd twdt w d dw w wdt wd tw dtdw wd
dwt d tagfdygafdgkdwlrn

fake wtdwt dwt wd twdt w d dw w wdt wd tw dtdw wd
dwt d tagfdygafdgkdwlrn



fake wtdwt dwt wd twdt w d dw w wdt wd tw dtdw wd
dwt d tagfdygafdgkdwlrn


fake wtdwt dwt wd twdt w d dw w wdt wd tw dtdw wd dwt
d tagfdygafdgkdwlrn fake wtdwt dwt wd twdt w d dw w wdt
wd tw dtdw wd dwt d tagfdygafdgkdwlrn fake wtdwt dwt wd
twdt w d dw w wdt wd tw dtdw wd dwt d tagfdygafdgkdwlrn
fake wtdwt dwt wd twdt w d dw w wdt wd tw dtdw wd dwt
d tagfdygafdgkdwlrn fake wtdwt dwt wd twdt w d dw w wdt
wd tw dtdw wd dwt d tagfdygafdgkdwlrn fake wtdwt dwt wd

REFERENCES


- [1] C. Jacobi, W. v. Atteveldt, and K. Welbers, "Quantitative analysis of large amounts of journalistic texts using topic modelling," *Digital Journalism*, vol. 4, no. 1, pp. 89–106, Jan. 2, 2016, ISSN: 2167-0811. DOI: 10.1080/21670811.2015.1093271.
- [2] P. Waila, V. K. Singh, and M. K. Singh, "Blog text analysis using topic modeling, named entity recognition and sentiment classifier combine," in *2013 International Conference on Advances in Computing, Communications and Informatics (ICACCI)*, Mysore: IEEE, Aug. 2013, pp. 1166–1171, ISBN: 978-1-4673-6217-7 978-1-4799-2432-5 978-1-4799-2659-6. DOI: 10.1109/ICACCI.2013.6637342.
- [3] J. Guo, P. Zhang, J. Tan, and L. Guo, "Mining hot topics from twitter streams," *Procedia Computer Science*, vol. 9, pp. 2008–2011, 2012, ISSN: 18770509. DOI: 10.1016/j.procs.2012.04.224.
- [4] J. Boyd-Graber, D. Mimno, and D. Newman, "Care and feeding of topic models: Problems, diagnostics, and improvements," *Handbook of Mixed Membership Models and Their Applications*, p. 30,



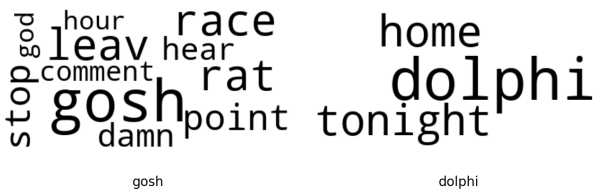
(a) Male




(b) Female



(c) 20 or younger



(d) Over 20



(e) Everyone

Fig. 4. Topics mined by TF-IDF. (a) male (b) female (c) 20 or younger (d) over 20 (e) everyone

twtdt w d dw w wdt wd tw dtdw wd dwt d tagfdygafdgkdwlrn
fake wtdwt dwt wd twdt w d dw w wdt wd tw dtdw wd dwt
d tagfdygafdgkdwlrn

fake wtdwt dwt wd twdt w d dw w wdt wd tw dtdw wd
dwt d tagfdygafdgkdwlrn

fake wtdwt dwt wd twdt w d dw w wdt wd tw dtdw wd
dwt d tagfdygafdgkdwlrn fake wtdwt dwt wd twdt w d dw w
wdt wd tw dtdw wd dwt d tagfdygafdgkdwlrn

fake wtdwt dwt wd twdt w d dw w wdt wd tw dtdw wd

APPENDIX

SOURCE CODE IN PYTHON

A. Code for topic mining

```
1 #!/usr/bin/env python
2 # -*- coding: utf-8 -*-
3
4 import sys
5 import os.path
6 from glob import glob
7 from tqdm import tqdm
8 import pickle
9 import json
10 from datetime import date
11 import pprint
12 pp = pprint.PrettyPrinter(indent=2)
13
14 import random
15 import itertools
16 from collections import namedtuple, Counter, OrderedDict, defaultdict
17 import heapq
18 from operator import itemgetter
19 import re
20 from bs4 import BeautifulSoup
21 import numpy as np
22
23 from spellchecker import SpellChecker
24 import nltk
25 from nltk.corpus import stopwords
26 from nltk.corpus import wordnet
27 from nltk.stem import PorterStemmer, LancasterStemmer, WordNetLemmatizer
28
29 NUM_SAMPLES = None
30 _DEBUG = False
31 NUM_SAMPLES = 10000
32 NUM_SAMPLES = 5000
33
34
35 STOPWORDS = set(stopwords.words("english"))
36 # Add more stopwords manually
37 with open('stopwords1.txt') as f:
38     STOPWORDS.update(w.strip().lower() for w in f)
39 STOPWORDS.update(['i\'m', 'dont', '\t', '\m', '\s', '\re', '\ve',
40     'haha', 'hah', 'wow', 'hehe', 'heh',
41     'ah', 'ahh', 'hm', 'hmm', 'urllink', 'ok', 'hey', 'yay', 'yeah'])
42
43 #####
44 #           Utility functions           #
45 #####
46
47 def len2d(iter2d):
48     return sum(len(d) for d in iter2d)
49
50 def list2d(iter2d):
51     return [[x for x in inner] for inner in iter2d]
52
53 def flatten2d(list2d):
54     return itertools.chain.from_iterable(list2d)
55
56 def flatten3d(list3d):
57     return itertools.chain.from_iterable(flatten2d(list3d))
58
59 def mapbar(f, seq, desc):
60     for e in tqdm(seq, desc):
61         yield f(e)
62
```

```

63 def map2d(f, docs):
64     with tqdm(total=len2d(docs)) as pbar:
65         def _helper(sent):
66             pbar.update(1)
67             return f(sent)
68
69         return [list(map(_helper, doc)) for doc in docs]
70
71 def map3d(f, docs):
72     with tqdm(total=len2d(docs)) as pbar:
73         def _helper(sent):
74             pbar.update(1)
75             return [f(word) for word in sent]
76
77         return [list(map(_helper, doc)) for doc in docs]
78
79 def foreach3d(f, docs):
80     with tqdm(total=len2d(docs)) as pbar:
81         for doc in docs:
82             for sent in doc:
83                 for word in sent:
84                     f(word)
85             pbar.update(1)
86
87 def foreach2d(f, docs):
88     with tqdm(total=len2d(docs)) as pbar:
89         for doc in docs:
90             for sent in doc:
91                 f(sent)
92             pbar.update(1)
93
94 def filter3d(f, docs):
95     ret = []
96     with tqdm(total=len2d(docs)) as pbar:
97         def _helper_doc(doc):
98             for sent in doc:
99                 pbar.update(1)
100                 out = [word for word in sent if f(word)]
101                 if len(out) > 0:
102                     yield out
103
104         for doc in docs:
105             out = list(_helper_doc(doc))
106             ret.append(out)
107     return ret
108
109 def load_pkl(fpath):
110     print('load dataset from cached pickle file ' + fpath)
111     with open(fpath, 'rb') as f:
112         dataset = pickle.load(f)
113     return dataset
114
115 def save_pkl(obj, fpath):
116     with open(fpath, 'wb') as f:
117         print('save dataset to pickle file ' + fpath)
118         pickle.dump(obj, f)
119
120 def save_json(obj, fpath, indent=2):
121     with open(fpath, 'w', encoding="utf8") as f:
122         print('save dataset to json file ' + fpath)
123         json.dump(obj, f, indent=indent)
124
125 #####
126 #           Codes for data reading & transformation           #
127 #####
128
129 Record = namedtuple('Record', ['meta', 'posts'])

```

```

130 Post = namedtuple('Post', ['date', 'text'])
131 MetaData = namedtuple('MetaData', ['id', 'gender', 'age', 'category', 'zodiac'])
132
133 def parse_meta_data(meta_data_str):
134     arr = meta_data_str.lower().strip().split('.')
135     return MetaData(arr[0], arr[1], int(arr[2]), arr[3], arr[4])
136
137 def read_blog_file(fpath):
138     try:
139         with open(fpath, encoding='utf-8', errors='ignore') as f:
140             soup = BeautifulSoup(f.read(), "xml")
141             blog = soup.Blog
142     except ParseError:
143         print('Error: invalid xml file {}'.format(fpath))
144         raise
145     return []
146
147     posts = []
148     state = 'date'
149     for c in blog.find_all(recursive=False):
150         if c.name != state:
151             print('Warning: inconsistent format in file {}'.format(fpath))
152         if state == 'date':
153             try:
154                 date_str = c.text.strip()
155                 date = date_str
156             except ValueError:
157                 print('Warning: invalid date {} in file {}' \
158                       .format(c.text, fpath))
159             state = 'post'
160         else:
161             text = c.text.strip()
162             state = 'date'
163             posts.append(Post(date, text))
164     posts.sort(key=lambda p: p.date)
165     return posts
166
167 def read_blogs(path, force=False, cache_file='blogs.pkl'):
168     if not force and cache_file is not None and os.path.exists(cache_file):
169         return load_pkl(cache_file)
170
171     dataset = read_blogs_xml(path)
172
173     # save to pickle file for fast loading next time
174     if cache_file is not None:
175         save_pkl(dataset, cache_file)
176
177     return dataset
178
179 def read_blogs_xml(path):
180     print('reading all data files from directory {} ...'.format(path))
181     dataset = []
182
183     if _DEBUG: # use small files for fast debugging
184         files = [os.path.join(path, fname) for fname in ['3998465.male.17.indUnk.Gemini.xml',
185                                                         '3949642.male.25.indUnk.Leo.xml', '3924311.male.27.HumanResources.Gemini.xml']]
186         files = random.sample(list(glob(os.path.join(path, '*'))), 100)
187     elif NUM_SAMPLES is None:
188         files = glob(os.path.join(path, '*'))
189     else:
190         files = random.sample(list(glob(os.path.join(path, '*'))), NUM_SAMPLES)
191
192     for fpath in tqdm(files):
193         fname = os.path.basename(fpath)
194         meta_data = parse_meta_data(fname)
195         posts = read_blog_file(fpath)
196         rec = Record(meta_data, posts)

```



```

197         dataset.append(rec)
198     return dataset
199
200 #####
201 #               Codes for topic mining               #
202 #####
203
204 punct_re = re.compile(r'([\.\!?,;:~])(?=[a-zA-Z])') # add space between a punctuation and a word
205 # replace two or more consecutive single quotes to a double quote
206 # e.g. ' ' -> " " ' ' -> " "
207 quotes_re = re.compile(r"['\"]{2,}")
208 def preprocess(text):
209     out = punct_re.sub(r'\1 ', text)
210     out = quotes_re.sub(r'"', out)
211     out = remove_invalid(out)
212     return out
213
214 leading_quote_re = re.compile(r'[\.\!?,;:~]*&^%#|!|-]+([a-zA-Z].*)')
215 def clean_word(word):
216     if word in ("ve", "re", "s", "t", "ll", "m", "d", "", "'"):
217         return word
218     word = leading_quote_re.sub(r'\1', word)
219     return word.strip()
220
221 def tokenise(dataset):
222     """
223     consider all the blogs from one person as a document
224
225     Returns
226     -----
227     docs: list of list of list
228         a list of documents, each of which is a list of sentences,
229         each of which is a list of words.
230     """
231
232     print('tokenising the text dataset...')
233     docs = []
234     with tqdm(total=sum(len(rec.posts) for rec in dataset)) as pbar:
235         for rec in dataset:
236             doc = []
237             for post in rec.posts:
238                 for sent_str in nltk.sent_tokenize(post.text):
239                     sent_str = preprocess(sent_str)
240                     sent = [clean_word(w) for w in nltk.word_tokenize(sent_str)]
241                     sent = [w for w in sent if w != '']
242                     doc.append(sent)
243             pbar.update(1)
244             docs.append(doc)
245
246     return docs
247
248 def calc_vocab(docs):
249     """Calculate the vocabulary (set of distinct words) from a collection
250     of documents.
251     """
252
253     print('calculating the vocabulary...')
254     vocab = set()
255
256     def _helper(sent):
257         vocab.update(sent)
258
259     foreach2d(_helper, docs)
260     return sorted(vocab)
261
262 def calc_pos_tags(docs):
263     print('POS tagging...')

```

```

264 def _f(sent):
265     try:
266         return nltk.pos_tag(sent)
267     except IndexError:
268         print('error sentence: {}'.format(sent))
269         raise
270 tagged_docs = map2d(_f, docs)
271 return tagged_docs
272
273 pattern = re.compile(r'([^\.] )\1{2,}')
274 pattern_ellipse = re.compile(r'\.{4,}')
275 invalid_chars = re.compile(r'[*\^#]')
276 def remove_invalid(text):
277     '''Basic cleaning of words, including:
278
279     1. rip off characters repeated more than twice as English words have a max
280     of two repeated characters.
281     2. remove characters which are not part of English words
282     '''
283
284     text = invalid_chars.sub(' ', text)
285     text = pattern.sub(r'\1\1', text)
286     text = pattern_ellipse.sub('...', text)
287     return text.strip()
288
289 def remove_invalid_all(docs):
290     print('reduce lengthily repeated characters...')
291     return filter3d(lambda w: len(w) > 0, map3d(remove_invalid, docs))
292
293 spell = SpellChecker()
294
295 def correct_spelling(word):
296     if not wordnet.synsets(word) and not word in STOPWORDS:
297         return spell.correction(word)
298     else:
299         return word
300
301 def correct_spelling_all(docs):
302     print('running spelling correction...')
303     return map3d(correct_spelling, docs)
304
305 def remove_stopwords(docs):
306     print('removing stopwords...')
307     return filter3d(lambda wp: wp[0].lower() not in STOPWORDS, docs)
308
309 lemmatizer = WordNetLemmatizer()
310 porter = PorterStemmer()
311 lancaster = LancasterStemmer()
312 def stem_word(word):
313     return porter.stem(lemmatizer.lemmatize(word))
314
315 def do_stemming(docs):
316     print('stemming or lemmatising words...')
317     return map3d(lambda wp: (stem_word(wp[0]), wp[1]), docs)
318
319 def calc_ne_all(docs):
320     print('extracting named entities...')
321     def _calc_ne(sent):
322         ne = []
323         for chunk in nltk.ne_chunk(sent):
324             if hasattr(chunk, 'label'):
325                 ne.append((' '.join(c[0] for c in chunk), chunk.label()))
326         return ne
327     return map2d(_calc_ne, docs)
328
329
330 def calc_df(docs):

```

```

331 df = defaultdict(lambda: 0)
332 for doc in docs:
333     for w in set(doc):
334         df[w] += 1
335 return df
336
337 def calc_tfidf(docs):
338     '''The original TF-IDF is a document-wise score. This function will
339     calculate the average TF-IDF on whole dataset as an overall scoring.
340     '''
341     tf_idf = defaultdict(lambda: 0)
342     df = calc_df(docs)
343     num_docs = len(docs)
344     for doc in docs:
345         counter = Counter(doc)
346         num_words = len(doc)
347         for token in set(doc):
348             tf = counter[token] / num_words
349             df_i = df[token]
350             idf = np.log(num_docs / df_i)
351             tf_idf[token] += tf * idf
352
353     for token in tf_idf:
354         tf_idf[token] /= df[token]
355
356     return tf_idf
357
358 def get_top_topics(named_entities, n=5, method='tf'):
359     print('calculating most popular topics by ' + method + '...')
360     if method == 'tf':
361         ranks = nltk.FreqDist(w for w, t in flatten3d(named_entities))
362         print(ranks.most_common(50))
363         ranks = dict(ranks)
364     elif method == 'tfidf':
365         ranks = calc_tfidf([[w for w, t in flatten2d(doc)] for doc in named_entities])
366     ranks = [(k, v) for k, v in ranks.items()]
367     print('\n largest:', heapq.nlargest(200, ranks, key=itemgetter(1)))
368     topics = heapq.nlargest(n, ranks, key=itemgetter(1))
369     print('topics: ', topics)
370     return topics
371
372 def get_surroundings(words, docs, n=4):
373     '''expand the topic to be 2 verb/noun before and 2 verb/noun after the topic
374     '''
375
376     print('get surrounding 2 nouns/verbs for words {}'.format(words))
377
378     sur = {}
379     for w, c in words:
380         sur[w] = Counter()
381
382     # POS tags list for searching verbs/nouns
383
384     def _helper(sent):
385         sent_w = [w for w, p in sent]
386         for w, c in words:
387             try:
388                 idx = sent_w.index(w)
389             except ValueError:
390                 continue
391
392             after = 0
393             vicinity = [sent[i] for i in [idx-2, idx-1, idx+1, idx+2]
394                         if i >= 0 and i < len(sent)]
395             for (wi, pi) in vicinity:
396                 if pi.startswith('N') or pi.startswith('V'):
397                     sur[w][wi] += 1

```

```

398     foreach2d(_helper, docs)
399     ret = []
400     for w, c in words:
401         ret.append({'topic': w, 'score': c, 'keywords': sur[w].most_common(n)})
402     return ret
403
404
405 def calc_intermediate_data(dataset):
406     docs = tokenise(dataset)
407     vocab = calc_vocab(docs)
408     print('Size of vocabulary: {}'.format(len(vocab)))
409     print(vocab[1:2000:2])
410     print(vocab[1:100000:100])
411
412
413     tagged_docs = calc_pos_tags(docs)
414     docs = vocab = None
415
416     named_entities = calc_ne_all(tagged_docs)
417
418     # Remove stopwords after POS tagging and NER finished
419     tagged_docs = remove_stopwords(tagged_docs)
420     named_entities = remove_stopwords(named_entities)
421
422     tagged_docs = do_stemming(tagged_docs)
423     named_entities = do_stemming(named_entities)
424     return tagged_docs, named_entities
425
426
427 def mine_topics(dataset, intermediate_data, group='all'):
428     print('-' * 80)
429     print('mining most popular topics for group ' + group)
430     print('-' * 80)
431     tagged_docs, named_entities = intermediate_data
432
433     if group != 'all':
434         if group == 'male' or group == 'female':
435             idx = [i for i, rec in enumerate(dataset) if rec.meta.gender == group]
436         elif group == '<=20':
437             idx = [i for i, rec in enumerate(dataset) if rec.meta.age <= 20]
438         elif group == '>20':
439             idx = [i for i, rec in enumerate(dataset) if rec.meta.age > 20]
440         else:
441             raise NotImplementedError()
442         tagged_docs = [tagged_docs[i] for i in idx]
443         named_entities = [named_entities[i] for i in idx]
444
445     print('selected docs: {}, {}'.format(len(tagged_docs), len(named_entities)))
446
447     ret = {}
448     num_keywords = 200
449     print('----- result from TFIDF -----')
450     topics = get_top_topics(named_entities, n=50, method='tfidf')
451     keywords = get_surroundings(topics, tagged_docs, n=num_keywords)
452     ret['tfidf'] = keywords
453
454     print('----- result from TF -----')
455     topics = get_top_topics(named_entities, n=50, method='tf')
456     keywords = get_surroundings(topics, tagged_docs, n=num_keywords)
457     ret['tf'] = keywords
458     return ret
459
460 def main_intermediate():
461     if not _DEBUG and NUM_SAMPLES is None:
462         dataset = read_blogs('blogs')
463     else:
464         dataset = read_blogs('blogs', cache_file=None)

```

```

465 intermediate_data = calc_intermediate_data(dataset)
466 save_pkl(intermediate_data, 'intermediate_data.pkl')
467 return dataset, intermediate_data
468
469 def main_mine_topics(dataset=None, intermediate_data=None):
470     if dataset is None:
471         dataset = load_pkl('blogs.pkl')
472     if intermediate_data is None:
473         intermediate_data = load_pkl('intermediate_data.pkl')
474
475     topics = {}
476     topics['male'] = mine_topics(dataset, intermediate_data, group='male')
477     topics['female'] = mine_topics(dataset, intermediate_data, group='female')
478     topics['less_or_20'] = mine_topics(dataset, intermediate_data, group='<=20')
479     topics['over_20'] = mine_topics(dataset, intermediate_data, group='>20')
480     topics['all'] = mine_topics(dataset, intermediate_data, group='all')
481     if _DEBUG:
482         suffix = 'debug'
483     else:
484         suffix = date.today().strftime('%Y%m%d')
485         if NUM_SAMPLES > 0:
486             suffix += '-' + str(NUM_SAMPLES)
487
488     save_json(topics, 'topics-{}.json'.format(suffix))
489
490 def main():
491     if len(sys.argv) <= 1:
492         phases = [1, 2]
493     else:
494         phases = [int(i) for i in sys.argv[1].split(',')]
495
496     dataset = intermediate_data = None
497     for ph in phases:
498         if ph == 1:
499             dataset, intermediate_data = main_intermediate()
500         elif ph == 2:
501             main_mine_topics(dataset, intermediate_data)
502
503 if __name__ == '__main__':
504     main()
505

```

B. Code for analysis, evaluation and visualisation

```

1  #!/usr/bin/env python
2  # -*- coding: utf-8 -*-
3
4  import sys
5  import json
6  import numpy as np
7  import pandas as pd
8  from as2 import load_pkl, Record, MetaData, Post
9  import matplotlib.pyplot as plt
10 from wordcloud import WordCloud
11
12 MAX_FONT_SIZE = 80
13
14 def show_summary(dataset):
15     '''This function describes the summary of dataset or human inspection.
16     It's not necessary for the mining process.
17
18     Parameters
19     -----
20     dataset : list of Record
21         The blog dataset
22     '''
23

```

```

24 df = pd.DataFrame([d.meta for d in dataset])
25 df['blog_count'] = [len(d.posts) for d in dataset]
26 df['char_count'] = [sum(len(p.text) for p in d.posts) for d in dataset]
27
28 print(df.describe(include='all'))
29 print('{} possible values for "gender": {}'.format(
30     len(df.gender.unique()), ', '.join(sorted(df.gender.unique()))))
31 print('{} possible values for category: {}'.format(
32     len(df.category.unique()), ', '.join(sorted(df.category.unique()))))
33 print('{} possible values for zodiac: {}'.format(
34     len(df.zodiac.unique()), ', '.join(sorted(df.zodiac.unique()))))
35
36 plt.rcParams.update({'font.size': 20})
37 df['gender'].value_counts().plot(kind='bar')
38 plt.xticks(rotation=0)
39 plt.gcf().tight_layout()
40 plt.savefig('img/show-gender.png')
41
42 plt.rcParams.update({'font.size': 10})
43 plt.clf()
44 df['category'].value_counts().plot(kind='bar')
45 plt.gcf().tight_layout()
46 plt.savefig('img/show-category.png')
47
48 plt.rcParams.update({'font.size': 18})
49 plt.clf()
50 df['zodiac'].value_counts().plot(kind='bar')
51 plt.xticks(rotation=90)
52 plt.gcf().tight_layout()
53 plt.savefig('img/show-zodiac.png')
54
55 plt.rcParams.update({'font.size': 20})
56 plt.clf()
57 age = df['age']
58 df['age'].hist(bins=20)
59 plt.gcf().tight_layout()
60 plt.savefig('img/show-age.png')
61
62 plt.clf()
63 cnt = df['blog_count']
64 logbins = np.logspace(np.log10(cnt.min()), np.log10(cnt.max()), 20)
65 cnt.hist(bins=logbins)
66 plt.xscale('log')
67 plt.gcf().tight_layout()
68 plt.savefig('img/show-blog-count.png')
69
70 plt.clf()
71 cnt = df['char_count']
72 logbins = np.logspace(np.log10(cnt.min()), np.log10(cnt.max()), 20)
73 cnt.hist(bins=logbins)
74 plt.xscale('log')
75 plt.gcf().tight_layout()
76 plt.savefig('img/show-char-count.png')
77
78 plt.clf()
79 df['gender_age'] = [g + '\n' + ('<=20' if a <= 20 else '>20') \
80     for (g, a) in zip(df['gender'], df['age'])]
81 df['gender_age'].value_counts()[[2, 3, 1, 0]].plot(kind='bar')
82 plt.xticks(rotation=0)
83 plt.gcf().tight_layout()
84 plt.savefig('img/show-gender-age.png')
85
86
87 def color_black(word, *args, **kwargs):
88     return '#000000'
89
90 def grey_color_func(word, font_size, position, orientation, random_state=None, **kwargs):

```

```

91     return 'hsl(0, 0%, {:d}%)'.format((MAX_FONT_SIZE - font_size) // (MAX_FONT_SIZE * 1))
92
93 def eval_topics(fpath, method='tf', top_k=2, num_words_in_topic=20):
94     with open(fpath, encoding='utf8') as f:
95         result = json.load(f)
96
97         for group, topics2 in result.items():
98             topics = topics2[method]
99             for i, topic in enumerate(topics[:top_k]):
100                 topic_name = topic['topic']
101                 words = {}
102                 words.update(tuple(kw for kw in topic['keywords'][:num_words_in_topic+1]))
103                 if method == 'tf':
104                     words[topic_name] = topic['score']
105                 else:
106                     words[topic_name] = topic['keywords'][0][1] * 2 # fake frequency for display
107
108                 print('topic: ', topic_name, 'number of keywords:', len(topic['keywords']))
109                 wc = WordCloud(background_color="white",
110                               max_font_size=80,
111                               max_words=num_words_in_topic+1,
112                               color_func=grey_color_func)
113                 wc.generate_from_frequencies(words)
114
115                 plt.clf()
116                 plt.imshow(wc, interpolation="bilinear")
117                 plt.axis("off")
118                 plt.title(topic_name, y=-0.25, fontsize=20)
119                 plt.gcf().tight_layout()
120                 fig_path = 'img/{}-{}-{}.png'.format(group, method, i+1, topic['topic'])
121                 print('drawing ' + fig_path)
122                 plt.savefig(fig_path)
123
124 def main():
125     cmd = sys.argv[1]
126     if cmd == 'show':
127         show_summary(load_pkl('blogs.pkl'))
128     elif cmd == 'eval':
129         fpath = sys.argv[2]
130         eval_topics(fpath, top_k=2)
131         eval_topics(fpath, top_k=2, method='tfidf')
132
133 if __name__ == '__main__':
134     main()

```