

Popular Topic Mining from Blogs

Stone Fang (Student ID: 19049045)
Computers and Information Sciences
Auckland University of Technology
Auckland, New Zealand
fkn7060@autuni.ac.nz

I. OVERVIEW

People's concerns and opinions are important reference for innovations of new products or services. However, accomplishing such task by humans is expensive, time-consuming and difficult to scale. As a response, a number of individuals and organisations are leveraging text mining technologies to mining meaningful information from large volume of text such as news media [1]. Among a variety of studies and applications, topic modelling is an important method to extract hot topics which reflects public attention and opinion from massive texts [1]–[3]. However, effective method of extracting useful information from text on the Internet remains an open challenge [3].

Evaluation of topics mined from text is another challenge, mostly due to the lack of ground truth because topic modelling is an unsupervised learning task [4].

The goal of this project is to mine most popular topics that people were discussing from blog posts by utilising various text mining algorithms and tools. Specifically, we will find two most popular topics for each group in the following demographics:

- Males
- Females
- People 20 years old or younger
- People older than 20
- Everyone

The remainder of this article is organised as follows. In section II related works on topic mining and evaluation will be reviewed. The methodology of topic mining are detailed in section III, while the results, analysis and evaluations are presented in section IV. The works of this article are summarised in section V and open issues and future works are discussed in section VI.

II. RELATED WORK

Jacobi, Atteveldt, and Welbers [1] conducted an in-depth study of how to apply topic modelling technologies on analysis of qualitative data in academic research.

III. RESEARCH DESIGN

In this section the solution will be described in detail. First an overview of the dataset is given, and then the algorithm of topic mining is detailed.

A. Data Description

The dataset contains 19,320 files in XML format, each containing articles of one person posted generally between 2001 and 2004. Metadata of the bloggers includes gender, age, category, and zodiac. In addition, the number of posts for each person are also counted. The result is summarised by Fig. 1, which is created by Python packages `Pandas` and `Matplotlib`. From this figure we can acquire some basic statistics of the dataset, including

- 1) Gender: data samples are quite evenly distributed over both genders.
- 2) Age: most bloggers are younger than 30, almost of them under 20. On the other hand, there are two gaps around 20 and 30 which may implies some missing data points in the dataset
- 3) Zodiac: The distribution over zodiac is reasonable even.
- 4) Number of posts: most bloggers published less than 100 posts, while the peak appears at 10, which implies people are most likely to write around 10 posts.
- 5) Category: the most frequent category is unknown, which is trivial, while the second frequent one is student, far more than other categories.

B. Topic Mining Algorithm

The general idea for mining popular topics used in this project is to find the most significant “things” mentioned in the overall dataset, as well as the closely related information.

The overall architecture of the algorithm is shown as Fig. 2, and the details of each step are described in the following subsections.

1) *Data Cleaning*: Before applying any text mining techniques, it is important to do basic data cleaning to improve data quality. In this step, a few operations for preprocessing will be carried out based on the observations of the dataset, with details as follows.

- **Problem**: At some place there is no whitespace between a punctuation and the word following it, which causes wrong tokenisation. Specifically, the punctuation might be tokenised with the following word as one token.
Solution: Add whitespace after a punctuation if a word immediately follows it.
- **Problem**: Two or more consecutive quote symbols may cause wrong tokenisation.
Solution: Replace two or more quotes as a double quote.

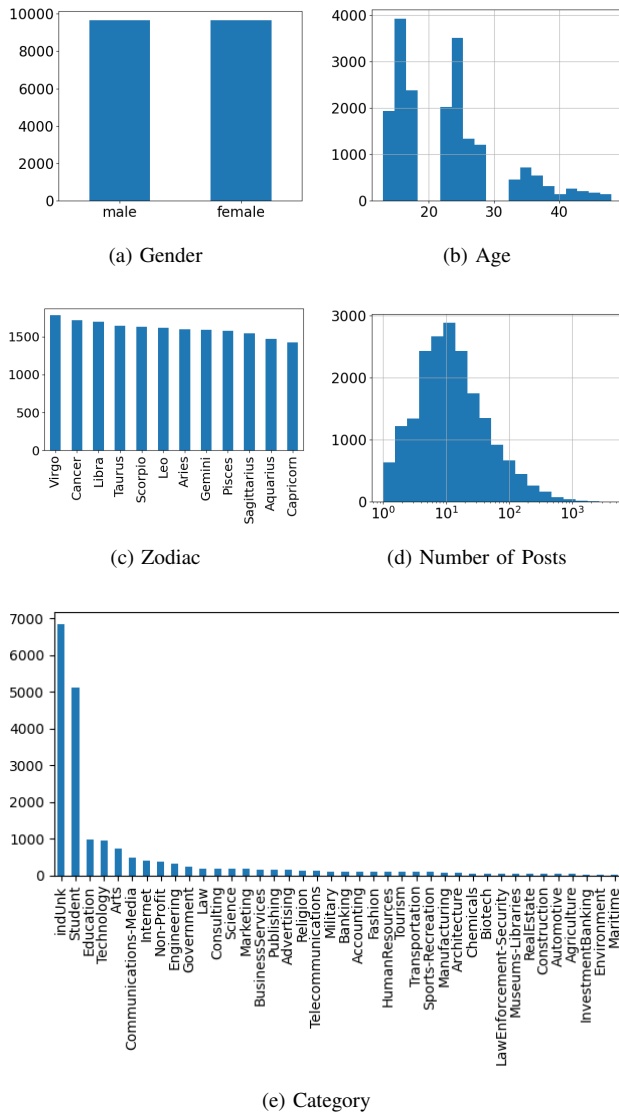


Fig. 1. Data Overview. Histogram over (a) gender (b) age (c) zodiac (d) number of posts (e) category

- **Problem:** The unicode quote may affect tokenisation and stopwords matching.
Solution: Replace unicode quote by ASCII quote.
- **Problem:** The unicode quote may affect tokenisation and stopwords matching.
Solution: Replace unicode quote by ASCII quote.
- **Problem:** Characters that are usually not part of normal English text may disturb tokenisation and POS tagging.
Solution: Remove invalid characters such as "*", ",", "#", and so on.
- **Problem:** Sometimes people repeat a certain letter in a word for emphasis, but it will result in wrong words and also increase the vocabulary size.
Solution: No English word has more than two consecutive appearances of the same letter, so three or more repetition of a letter is squeezed into two.

• Problem:

Solution:

These operations are implemented by regex matching and substitution, or simple text replacing. To use regex, the Python's `re` package are imported.

2) *Tokenisation*: Tokenisation is usually the first step of all text mining pipelines, which includes sentence and word tokenisation. Sentence tokenisation is to split the whole text into sentences, while word tokenisation splits a sentence into word or tokens. In this project we use `nltk` package to do such task. This package provides two functions `sent_tokenize()` and `word_tokenize()` for both tokenisation. A document is first tokenised into sentences, and then each sentence is tokenised into words. Finally, a document is represented as list of lists, as each sentence is a list of words.

3) *POS Tagging*: Part-Of-Speech (POS) tagging is the second step following tokenisation. In this step, each word is assigned by a POS tag. `nltk` provides a handy function `pos_tag()` to do this task. This function works on sentence level, and maps each word into a tuple which is the pair of word and POS tag.

4) *Entity Extraction*: In this project, a topic is defined as a "thing" or "object". Therefore, in order to find the topics, we need to find all "things" or "objects" first. There are a few options to do this task, among which two methods will be employed by this project: Named Entity Recognition (NER) and parsing.

a) *NER*: Named entities are ideal candidates of topics as they denote real-world objects. `nltk` provides a function `ne_chunk()` to extract entities from sentences. The input of the function should be a list of tokens with POS tags, which is another reason why POS tagging should be done in previous step. The return value of this function is a list of chunks, each of which is basically a list and may contain a `label` attribute if it is a recognised entity. The entity type can be acquired by `label()` and the entity itself should be acquired by joining all the elements of the chunk.

b) *Parsing*: Another way to extract objects is parsing by pre-defined patterns. For example, it is reasonable to treat definite nouns as objects according to the grammar.

5) *Stopwords Removal*: Stopwords are most common words which carries no significant meanings. Removing stopwords can reduce the size of data to be proceeded as well as increase the result accuracy. `nltk` provides an out-of-the-box stopwords collection, but the experiment shows that some common words carrying no meaning are not included in the list. In order to expand the stopwords list, more words are collected from website ¹.

Stopwords removal is conducted after POS tagging and entity extraction because these two steps are sequence model, which means their performance rely on word order. If stopwords are removed before them, we will get sentences which do not comply with English grammar. In addition, stopwords removal are carried out on tagged documents as well as entities

¹<https://gist.github.com/sebleier/554280>

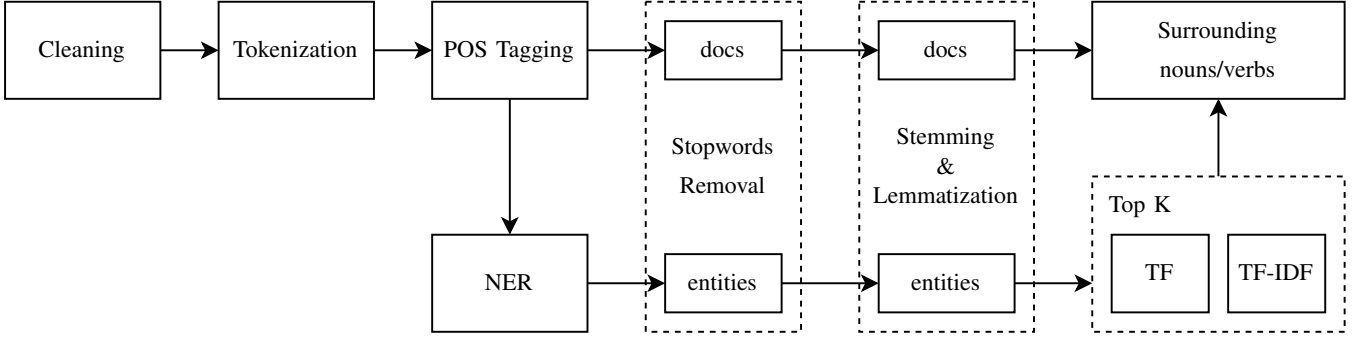


Fig. 2. Algorithm

extracted. Theoretically, stopwords cannot be entity, but errors will happen in any POS tagging and NER model. Therefore, trying to remove stopwords can reduce the error introduced in previous steps.

6) *Stemming and Lemmatisation*: Stemming and lemmatisation are both techniques for text normalisation, that is, convert an inflected word into its root form. However, stemming and lemmatisation work in different way. Stemming removes suffix or prefix from a word, returning a word stem which is not necessarily a word. On the other hand, lemmatisation always looks for the lemma from word variations with morphological analysis. For example, stemming against the third-person singular form “flies” returns “fli”, while lemmatisation returns “fly”. In this project, these two methods are combined together to reach the maximum extent of word normalisation.

`nltk` provides various stemming algorithms such as `PorterStemmer` and `LancasterStemmer`, and one lemmatisation algorithm `WordNetLemmatizer`. In the code we use `WordNetLemmatizer` followed by `PorterStemmer`.

7) *Word Count and TF-IDF*: After all “objects” have been extracted and normalised, the next step is to find most popular ones as the most dominant topics. Popularity can be defined in various ways, and in this project two approaches are used: word count and Term Frequency-Inverse Document Frequency (TF-IDF). In the first method, we simply count the appearances of each entity and get the most two frequent ones. In the second method, we calculate the TF-IDF value of each entity word, following the definition

$$\begin{aligned} \text{TF-IDF}(t_i, d_j) &= \text{TF}(t_i, d_j) \times \text{IDF}(t_i) \\ &= \text{TF}(t_i, d_j) \log \frac{N}{\text{DF}(t_i)} \end{aligned}$$

$\text{TF}(t_i, d_j)$ is the Term Frequency of term t_i in document d_j , which is computed by count of t_i in d_j divided by the total number of terms in d_j . $\text{DF}(t_i)$ is the Document Frequency, which is the number of documents that contains t_i . As we can see here, TF-IDF is a term-document-wise number so a term has different TF-IDF values in different documents. In order to rank all terms over the whole dataset, TF-IDF values of a term are averaged over all documents as the score of that term.

$$\text{score}(t_i) = \text{avg}_{d_j} \text{TF-IDF}(t_i, d_j)$$

Two different methods might return different results, which will be compared and analysed in section IV.

C. Evaluation

Evaluation of topics is challenging due to its nature of unsupervised learning. Among existing metrics, xx is chosen to evaluate the result of the methodology.

IV. RESULT, ANALYSIS, AND EVALUATION

A. Result

The results are presented by word cloud. Fig. 3 shows the topics mined by word count while Fig. 4 shows that by TF-IDF.

B. Analysis

C. Evaluation

V. CONCLUSION

This project has designed and implemented a complete solution to mine most popular topics from blogs. A variety of text mining technologies are employed and combined together to reach the goal. The results are compared and evaluated. Further and in-depth discussion is also provided.

VI. OPEN ISSUES AND FUTURE WORKS


There are still a few open issues remaining in the solution which can be improved by future work or changed if re-do this project.

fake wtdwt dwt wd twdt w d dw w wdt wd tw dtdw wd
dwt d tagfdygafdgkdwlrn gdwgd fake wtdwt dwt wd twdt w
d dw w wdt wd tw dtdw wd dwt d tagfdygafdgkdwlrn fake
wtdwt dwt wd twdt w d dw w wdt wd tw dtdw wd dwt d
tagfdygafdgkdwlrn


fake wtdwt dwt wd twdt w d dw w wdt wd tw dtdw wd
dwt d tagfdygafdgkdwlrn fake wtdwt dwt wd twdt w d dw w
wdt wd tw dtdw wd dwt d tagfdygafdgkdwlrn

fake wtdwt dwt wd twdt w d dw w wdt wd tw dtdw wd
dwt d tagfdygafdgkdwlrn


fake wtdwt dwt wd twdt w d dw w wdt wd tw dtdw wd
dwt d tagfdygafdgkdwlrn




(a) Male



iraq



god



bush

(a) Male



god



good

(b) Female




god



love

(b) Female



god



good

(c) 20 or younger




god



love

(c) 20 or younger



bush




god

(d) Over 20




god




bush

(d) Over 20




god



bush

(e) Everyone



god



love

(e) Everyone

Fig. 3. Topics mined by word count. (a) male (b) female (c) 20 or younger (d) over 20 (e) everyone

Fig. 4. Topics mined by TF-IDF. (a) male (b) female (c) 20 or younger (d) over 20 (e) everyone

fake wtdwt dwt wd twdt w d dw w wdt wd tw dtdw wd dwt d tagfdygafdgkdwlrn

fake wtdwt dwt wd twdt w d dw w wdt wd tw dtdw wd dwt d tagfdygafdgkdwlrn

fake wtdwt dwt wd twdt w d dw w wdt wd tw dtdw wd dwt d tagfdygafdgkdwlrn fake wtdwt dwt wd twdt w d dw w wdt wd tw dtdw wd dwt d tagfdygafdgkdwlrn fake wtdwt dwt wd twdt w d dw w wdt wd tw dtdw wd dwt d tagfdygafdgkdwlrn fake wtdwt dwt wd twdt w d dw w wdt wd tw dtdw wd dwt d tagfdygafdgkdwlrn

d tagfdygafdgkdwlrn fake wtdwt dwt wd twdt w d dw w wdt wd tw dtdw wd dwt d tagfdygafdgkdwlrn fake wtdwt dwt wd twdt w d dw w wdt wd tw dtdw wd dwt d tagfdygafdgkdwlrn fake wtdwt dwt wd twdt w d dw w wdt wd tw dtdw wd dwt d tagfdygafdgkdwlrn

fake wtdwt dwt wd twdt w d dw w wdt wd tw dtdw wd dwt d tagfdygafdgkdwlrn

fake wtdwt dwt wd twdt w d dw w wdt wd tw dtdw wd dwt d tagfdygafdgkdwlrn fake wtdwt dwt wd twdt w d dw w wdt wd tw dtdw wd dwt d tagfdygafdgkdwlrn

wdt wd tw dtdw wd dwt d tagfdygafdgkdwlrn
fake wtdwt dwt wd twdt w d dw w wdt wd tw dtdw wd
dwt d tagfdygafdgkdwlrn

REFERENCES

- [1] C. Jacobi, W. v. Atteveldt, and K. Welbers, “Quantitative analysis of large amounts of journalistic texts using topic modelling,” *Digital Journalism*, vol. 4, no. 1, pp. 89–106, Jan. 2, 2016, ISSN: 2167-0811. DOI: 10.1080/21670811.2015.1093271.
- [2] P. Waila, V. K. Singh, and M. K. Singh, “Blog text analysis using topic modeling, named entity recognition and sentiment classifier combine,” in *2013 International Conference on Advances in Computing, Communications and Informatics (ICACCI)*, Mysore: IEEE, Aug. 2013, pp. 1166–1171, ISBN: 978-1-4673-6217-7 978-1-4799-2432-5 978-1-4799-2659-6. DOI: 10.1109/ICACCI.2013.6637342.
- [3] J. Guo, P. Zhang, J. Tan, and L. Guo, “Mining hot topics from twitter streams,” *Procedia Computer Science*, vol. 9, pp. 2008–2011, 2012, ISSN: 18770509. DOI: 10.1016/j.procs.2012.04.224.
- [4] J. Boyd-Graber, D. Mimno, and D. Newman, “Care and feeding of topic models: Problems, diagnostics, and improvements,” *Handbook of Mixed Membership Models and Their Applications*, p. 30,

APPENDIX

SOURCE CODE IN PYTHON

A. Code for topic mining

```
1 #!/usr/bin/env python
2 # -*- coding: utf-8 -*-
3
4 import sys
5 import os.path
6 from glob import glob
7 from tqdm import tqdm
8 import pickle
9 import json
10 from datetime import date
11 import pprint
12 pp = pprint.PrettyPrinter(indent=2)
13
14 import random
15 import itertools
16 from collections import namedtuple, Counter, OrderedDict, defaultdict
17 import heapq
18 from operator import itemgetter
19 import re
20 from bs4 import BeautifulSoup
21 import numpy as np
22
23 from spellchecker import SpellChecker
24 import nltk
25 from nltk.corpus import stopwords
26 from nltk.corpus import wordnet
27 from nltk.stem import PorterStemmer, LancasterStemmer, WordNetLemmatizer
28
29 NUM_SAMPLES = None
30 _DEBUG = False
31
32 _DEBUG = True
33
34 STOPWORDS = set(stopwords.words("english"))
35 # Add more stopwords manually
36 with open('stopwords1.txt') as f:
37     STOPWORDS.update(w.strip().lower() for w in f)
38 STOPWORDS.update(['i\'m', 'dont', '\t', '\m', '\s', '\re', '\ve',
39                  'haha', 'hah', 'wow', 'hehe', 'heh',
40                  'ah', 'ahh', 'hm', 'hmm', 'urllink', 'ok', 'hey', 'yay', 'yeah'])
41
42 #####
43 #                               Utility functions                               #
44 #####
45
46 def len2d(iter2d):
47     return sum(len(d) for d in iter2d)
48
49 def list2d(iter2d):
50     return [[x for x in inner] for inner in iter2d]
51
52 def flatten2d(list2d):
53     return itertools.chain.from_iterable(list2d)
54
55 def flatten3d(list3d):
56     return itertools.chain.from_iterable(flatten2d(list3d))
57
58 def mapbar(f, seq, desc):
59     for e in tqdm(seq, desc):
60         yield f(e)
61
62 def map2d(f, docs):
```

```

63     with tqdm(total=len2d(docs)) as pbar:
64         def _helper(sent):
65             pbar.update(1)
66             return f(sent)
67
68         return [list(map(_helper, doc)) for doc in docs]
69
70 def map3d(f, docs):
71     with tqdm(total=len2d(docs)) as pbar:
72         def _helper(sent):
73             pbar.update(1)
74             return [f(word) for word in sent]
75
76         return [list(map(_helper, doc)) for doc in docs]
77
78 def foreach3d(f, docs):
79     with tqdm(total=len2d(docs)) as pbar:
80         for doc in docs:
81             for sent in doc:
82                 for word in sent:
83                     f(word)
84             pbar.update(1)
85
86 def foreach2d(f, docs):
87     with tqdm(total=len2d(docs)) as pbar:
88         for doc in docs:
89             for sent in doc:
90                 f(sent)
91             pbar.update(1)
92
93 def filter3d(f, docs):
94     ret = []
95     with tqdm(total=len2d(docs)) as pbar:
96         def _helper_doc(doc):
97             for sent in doc:
98                 pbar.update(1)
99                 out = [word for word in sent if f(word)]
100                 if len(out) > 0:
101                     yield out
102
103         for doc in docs:
104             out = list(_helper_doc(doc))
105             ret.append(out)
106     return ret
107
108 def load_pkl(fpath):
109     print('load dataset from cached pickle file ' + fpath)
110     with open(fpath, 'rb') as f:
111         dataset = pickle.load(f)
112     return dataset
113
114 def save_pkl(obj, fpath):
115     with open(fpath, 'wb') as f:
116         print('save dataset to pickle file ' + fpath)
117         pickle.dump(obj, f)
118
119 def save_json(obj, fpath, indent=2):
120     with open(fpath, 'w', encoding="utf8") as f:
121         print('save dataset to json file ' + fpath)
122         json.dump(obj, f, indent=indent)
123
124 #####
125 #           Codes for data reading & transformation           #
126 #####
127
128 Record = namedtuple('Record', ['meta', 'posts'])
129 Post = namedtuple('Post', ['date', 'text'])

```

```

130 MetaData = namedtuple('MetaData', ['id', 'gender', 'age', 'category', 'zodiac'])
131
132 def parse_meta_data(meta_data_str):
133     arr = meta_data_str.lower().strip().split('.')
134     return MetaData(arr[0], arr[1], int(arr[2]), arr[3], arr[4])
135
136 def read_blog_file(fpath):
137     try:
138         with open(fpath, encoding='utf-8', errors='ignore') as f:
139             soup = BeautifulSoup(f.read(), "xml")
140             blog = soup.Blog
141     except ParseError:
142         print('Error: invalid xml file {}'.format(fpath))
143         raise
144     return []
145
146 posts = []
147 state = 'date'
148 for c in blog.find_all(recursive=False):
149     if c.name != state:
150         print('Warning: inconsistent format in file {}'.format(fpath))
151     if state == 'date':
152         try:
153             date_str = c.text.strip()
154             date = date_str
155         except ValueError:
156             print('Warning: invalid date {} in file {}' \
157                   .format(c.text, fpath))
158             state = 'post'
159     else:
160         text = c.text.strip()
161         state = 'date'
162         posts.append(Post(date, text))
163 posts.sort(key=lambda p: p.date)
164 return posts
165
166 def read_blogs(path, force=False, cache_file='blogs.pkl'):
167     if not force and cache_file is not None and os.path.exists(cache_file):
168         return load_pkl(cache_file)
169
170     dataset = read_blogs_xml(path)
171
172     # save to pickle file for fast loading next time
173     if cache_file is not None:
174         save_pkl(dataset, cache_file)
175
176     return dataset
177
178 def read_blogs_xml(path):
179     print('reading all data files from directory {} ...'.format(path))
180     dataset = []
181
182     if _DEBUG: # use small files for fast debugging
183         files = [os.path.join(path, fname) for fname in ['3998465.male.17.indUnk.Gemini.xml',
184                                                         '3949642.male.25.indUnk.Leo.xml', '3924311.male.27.HumanResources.Gemini.xml']]
185         files = random.sample(list(glob(os.path.join(path, '*'))), 100)
186     elif NUM_SAMPLES is None:
187         files = glob(os.path.join(path, '*'))
188     else:
189         files = random.sample(list(glob(os.path.join(path, '*'))), NUM_SAMPLES)
190
191     for fpath in tqdm(files):
192         fname = os.path.basename(fpath)
193         meta_data = parse_meta_data(fname)
194         posts = read_blog_file(fpath)
195         rec = Record(meta_data, posts)
196         dataset.append(rec)

```



```

197     return dataset
198
199 #####
200 #               Codes for topic mining               #
201 #####
202
203 punct_re = re.compile(r'([\.\!?,;:;])(?=[a-zA-Z])') # add space between a punctuation and a word
204 # replace two or more consecutive single quotes to a double quote
205 #   e.g. ''' -> "        ''' -> "
206 quotes_re = re.compile(r"['\"]{2,}")
207 def preprocess(text):
208     print(text)
209     out = punct_re.sub(r'\1 ', text)
210     print(out)
211     out = quotes_re.sub(r'"', out)
212     print(out)
213     print(out)
214     out = remove_invalid(out)
215     print(out)
216     return out
217
218 leading_quote_re = re.compile(r'[\.\!?,;:;~*^%#|!|-]+([a-zA-Z].*)')
219 def clean_word(word):
220     if word in ("ve", "re", "s", "t", "ll", "m", "d", "n", "'"):
221         return word
222     word = leading_quote_re.sub(r'\1', word)
223     return word.strip()
224
225 def tokenise(dataset):
226     """
227     consider all the blogs from one person as a document
228
229     Returns
230     -----
231     docs: list of list of list
232         a list of documents, each of which is a list of sentences,
233         each of which is a list of words.
234     """
235
236     print('tokenising the text dataset...')
237     docs = []
238     with tqdm(total=sum(len(rec.posts) for rec in dataset)) as pbar:
239         for rec in dataset:
240             doc = []
241             for post in rec.posts:
242                 for sent_str in nltk.sent_tokenize(post.text):
243                     sent_str = preprocess(sent_str)
244                     sent = [clean_word(w) for w in nltk.word_tokenize(sent_str)]
245                     sent = [w for w in sent if w != '']
246                     doc.append(sent)
247             pbar.update(1)
248             docs.append(doc)
249
250     return docs
251
252 def calc_vocab(docs):
253     """Calculate the vocabulary (set of distinct words) from a collection
254     of documents.
255     """
256
257     print('calculating the vocabulary...')
258     vocab = set()
259
260     def _helper(sent):
261         vocab.update(sent)
262
263     foreach2d(_helper, docs)

```

```

264     return sorted(vocab)
265
266 def calc_pos_tags(docs):
267     print('POS tagging...')
268     def _f(sent):
269         try:
270             return nltk.pos_tag(sent)
271         except IndexError:
272             print('error sentence: {}'.format(sent))
273             raise
274     tagged_docs = map2d(_f, docs)
275     return tagged_docs
276
277 pattern = re.compile(r'([^\.])(\1{2,})')
278 pattern_ellipse = re.compile(r'\.{4,}')
279 invalid_chars = re.compile(r'[*^#]')
280 def remove_invalid(text):
281     '''Basic cleaning of words, including:
282
283     1. rip off characters repeated more than twice as English words have a max
284     of two repeated characters.
285     2. remove characters which are not part of English words
286     '''
287
288     print(text)
289     text = invalid_chars.sub(' ', text)
290     print(text)
291     text = pattern.sub(r'\1\1', text)
292     print(text)
293     text = pattern_ellipse.sub('...', text)
294     print(text)
295     return text.strip()
296
297 def remove_invalid_all(docs):
298     print('reduce lengthily repeated characters...')
299     return filter3d(lambda w: len(w) > 0, map3d(remove_invalid, docs))
300
301 spell = SpellChecker()
302
303 def correct_spelling(word):
304     if not wordnet.synsets(word) and not word in STOPWORDS:
305         return spell.correction(word)
306     else:
307         return word
308
309 def correct_spelling_all(docs):
310     print('running spelling correction...')
311     return map3d(correct_spelling, docs)
312
313 def remove_stopwords(docs):
314     print('removing stopwords...')
315     return filter3d(lambda wp: wp[0].lower() not in STOPWORDS, docs)
316
317 lemmatizer = WordNetLemmatizer()
318 porter = PorterStemmer()
319 lancaster = LancasterStemmer()
320 def stem_word(word):
321     return lemmatizer.lemmatize(word)
322
323 def do_stemming(docs):
324     print('stemming or lemmatising words...')
325     return map3d(lambda wp: (stem_word(wp[0]), wp[1]), docs)
326
327 def calc_ne_all(docs):
328     print('extracting named entities...')
329     def _calc_ne(sent):
330         ne = []

```

```

331         for chunk in nltk.ne_chunk(sent):
332             if hasattr(chunk, 'label'):
333                 ne.append((' '.join(c[0] for c in chunk), chunk.label()))
334         return ne
335     return map2d(_calc_ne, docs)
336
337
338 def calc_df(docs):
339     df = defaultdict(lambda: 0)
340     for doc in docs:
341         for w in set(doc):
342             df[w] += 1
343     return df
344
345 def calc_tfidf(docs):
346     '''The original TF-IDF is a document-wise score. This function will
347     calculate the average TF-IDF on whole dataset as an overall scoring.
348     '''
349     tf_idf = defaultdict(lambda: 0)
350     df = calc_df(docs)
351     num_docs = len(docs)
352     for doc in docs:
353         counter = Counter(doc)
354         num_words = len(doc)
355         for token in set(doc):
356             tf = counter[token] / num_words
357             df_i = df[token]
358             idf = np.log(num_docs / df_i)
359             tf_idf[token] += tf * idf
360
361     for token in tf_idf:
362         tf_idf[token] /= num_docs
363
364     return tf_idf
365
366 def get_top_topics(named_entities, n=5, method='tf'):
367     print('calculating most popular topics by ' + method + '...')
368     if method == 'tf':
369         ranks = nltk.FreqDist(w for w, t in flatten3d(named_entities))
370         print(ranks.most_common(50))
371         ranks = dict(ranks)
372     elif method == 'tfidf':
373         ranks = calc_tfidf([[w for w, t in flatten2d(doc)] for doc in named_entities])
374     ranks = [(k, v) for k, v in ranks.items()]
375     print('\n largest:', heapq.nlargest(200, ranks, key=itemgetter(1)))
376     topics = heapq.nlargest(n, ranks, key=itemgetter(1))
377     print('topics: ', topics)
378     return topics
379
380 def get_surroundings(words, docs, n=4):
381     '''expand the topic to be 2 verb/noun before and 2 verb/noun after the topic
382     '''
383
384     print('get surrounding 2 nouns/verbs for words {}'.format(words))
385
386     sur = {}
387     for w, c in words:
388         sur[w] = Counter()
389
390     # POS tags list for searching verbs/nouns
391
392     def _helper(sent):
393         sent_w = [w for w, p in sent]
394         for w, c in words:
395             try:
396                 idx = sent_w.index(w)
397                 except ValueError:

```

```

398         continue
399
400     after = 0
401     vicinity = [sent[i] for i in [idx-2, idx-1, idx+1, idx+2]
402                 if i >= 0 and i < len(sent)]
403     for (wi, pi) in vicinity:
404         if pi.startswith('N') or pi.startswith('V'):
405             sur[w][wi] += 1
406
407     foreach2d(_helper, docs)
408     ret = []
409     for w, c in words:
410         ret.append({'topic': w, 'score': c, 'keywords': sur[w].most_common(n)})
411     return ret
412
413 def calc_intermediate_data(dataset):
414     docs = tokenise(dataset)
415     vocab = calc_vocab(docs)
416     print('Size of vocabulary: {}'.format(len(vocab)))
417     print(vocab[1:2000:2])
418     print(vocab[1:100000:100])
419
420
421     tagged_docs = calc_pos_tags(docs)
422     docs = vocab = None
423
424     named_entities = calc_ne_all(tagged_docs)
425
426     # Remove stopwords after POS tagging and NER finished
427     tagged_docs = remove_stopwords(tagged_docs)
428     named_entities = remove_stopwords(named_entities)
429
430     tagged_docs = do_stemming(tagged_docs)
431     named_entities = do_stemming(named_entities)
432     return tagged_docs, named_entities
433
434
435 def mine_topics(dataset, intermediate_data, group='all'):
436     print('--' * 80)
437     print('mining most popular topics for group ' + group)
438     print('--' * 80)
439     tagged_docs, named_entities = intermediate_data
440
441     if group != 'all':
442         if group == 'male' or group == 'female':
443             idx = [i for i, rec in enumerate(dataset) if rec.meta.gender == group]
444         elif group == '<=20':
445             idx = [i for i, rec in enumerate(dataset) if rec.meta.age <= 20]
446         elif group == '>20':
447             idx = [i for i, rec in enumerate(dataset) if rec.meta.age > 20]
448         else:
449             raise NotImplementedError()
450         tagged_docs = [tagged_docs[i] for i in idx]
451         named_entities = [named_entities[i] for i in idx]
452
453     print('selected docs: {}, {}'.format(len(tagged_docs), len(named_entities)))
454
455     ret = {}
456     num_keywords = 200
457     print('----- result from TFIDF -----')
458     topics = get_top_topics(named_entities, n=50, method='tfidf')
459     keywords = get_surroundings(topics, tagged_docs, n=num_keywords)
460     ret['tfidf'] = keywords
461
462     print('----- result from TF -----')
463     topics = get_top_topics(named_entities, n=50, method='tf')
464     keywords = get_surroundings(topics, tagged_docs, n=num_keywords)

```

```

465     ret['tf'] = keywords
466     return ret
467
468 def main_intermediate():
469     if not _DEBUG and NUM_SAMPLES is None:
470         dataset = read_blogs('blogs')
471     else:
472         dataset = read_blogs('blogs', cache_file=None)
473
474     intermediate_data = calc_intermediate_data(dataset)
475     save_pkl(intermediate_data, 'intermediate_data.pkl')
476     return dataset, intermediate_data
477
478 def main_mine_topics(dataset=None, intermediate_data=None):
479     if dataset is None:
480         dataset = load_pkl('blogs.pkl')
481     if intermediate_data is None:
482         intermediate_data = load_pkl('intermediate_data.pkl')
483
484     topics = {}
485     topics['male'] = mine_topics(dataset, intermediate_data, group='male')
486     topics['female'] = mine_topics(dataset, intermediate_data, group='female')
487     topics['less_or_20'] = mine_topics(dataset, intermediate_data, group='<=20')
488     topics['over_20'] = mine_topics(dataset, intermediate_data, group='>20')
489     topics['all'] = mine_topics(dataset, intermediate_data, group='all')
490     if _DEBUG:
491         suffix = 'debug'
492     else:
493         suffix = date.today().strftime('%Y%m%d')
494         if NUM_SAMPLES > 0:
495             suffix += '-' + str(NUM_SAMPLES)
496
497     save_json(topics, 'topics-{}.json'.format(suffix))
498
499 def main():
500     if len(sys.argv) <= 1:
501         phases = [1, 2]
502     else:
503         phases = [int(i) for i in sys.argv[1].split(',')]
504
505     dataset = intermediate_data = None
506     for ph in phases:
507         if ph == 1:
508             dataset, intermediate_data = main_intermediate()
509         elif ph == 2:
510             main_mine_topics(dataset, intermediate_data)
511
512 if __name__ == '__main__':
513     main()

```

B. Code for analysis, evaluation and visualisation

```

1  #!/usr/bin/env python
2  # -*- coding: utf-8 -*-
3
4  import sys
5  import json
6  import numpy as np
7  import pandas as pd
8  from as2 import load_pkl, Record, MetaData, Post
9  import matplotlib.pyplot as plt
10 from wordcloud import WordCloud
11
12
13 def show_summary(dataset):
14     '''This function describes the summary of dataset or human inspection.
15     It's not necessary for the mining process.

```

```

16
17 Parameters
18 -----
19 dataset : list of Record
20           The blog dataset
21 '''
22
23 df = pd.DataFrame([d.meta for d in dataset])
24 df['blog_count'] = [len(d.posts) for d in dataset]
25 df['char_count'] = [sum(len(p.text) for p in d.posts) for d in dataset]
26
27 print(df.describe(include='all'))
28 print('{} possible values for "gender": {}'.format(
29     len(df.gender.unique()), ', '.join(sorted(df.gender.unique()))))
30 print('{} possible values for category: {}'.format(
31     len(df.category.unique()), ', '.join(sorted(df.category.unique()))))
32 print('{} possible values for zodiac: {}'.format(
33     len(df.zodiac.unique()), ', '.join(sorted(df.zodiac.unique()))))
34
35 plt.rcParams.update({'font.size': 20})
36 df['gender'].value_counts().plot(kind='bar')
37 plt.xticks(rotation=0)
38 plt.gcf().tight_layout()
39 plt.savefig('img/show-gender.png')
40
41 plt.rcParams.update({'font.size': 10})
42 plt.clf()
43 df['category'].value_counts().plot(kind='bar')
44 plt.gcf().tight_layout()
45 plt.savefig('img/show-category.png')
46
47 plt.rcParams.update({'font.size': 18})
48 plt.clf()
49 df['zodiac'].value_counts().plot(kind='bar')
50 plt.xticks(rotation=90)
51 plt.gcf().tight_layout()
52 plt.savefig('img/show-zodiac.png')
53
54 plt.rcParams.update({'font.size': 20})
55 plt.clf()
56 age = df['age']
57 df['age'].hist(bins=20)
58 plt.gcf().tight_layout()
59 plt.savefig('img/show-age.png')
60
61 plt.clf()
62 cnt = df['blog_count']
63 logbins = np.logspace(np.log10(cnt.min()), np.log10(cnt.max()), 20)
64 cnt.hist(bins=logbins)
65 plt.xscale('log')
66 plt.gcf().tight_layout()
67 plt.savefig('img/show-blog-count.png')
68
69 plt.clf()
70 cnt = df['char_count']
71 logbins = np.logspace(np.log10(cnt.min()), np.log10(cnt.max()), 20)
72 cnt.hist(bins=logbins)
73 plt.xscale('log')
74 plt.gcf().tight_layout()
75 plt.savefig('img/show-char-count.png')
76
77 def eval_topics(fpath, method='tf', top_k=2, num_words_in_topic=30):
78     with open(fpath, encoding='utf8') as f:
79         result = json.load(f)
80
81     for group, topics2 in result.items():
82         topics = topics2[method]

```

```

83     for i, topic in enumerate(topics[:top_k]):
84         topic_name = topic['topic']
85         words = {}
86         words.update(tuple(kw) for kw in topic['keywords'][:num_words_in_topic+1])
87         if method == 'tf':
88             words[topic_name] = topic['score']
89         else:
90             words[topic_name] = topic['keywords'][0][1] * 2 # fake frequency for display
91
92         print('topic: ', topic_name, 'number of keywords:', len(topic['keywords']))
93         wc = WordCloud(background_color="white", max_font_size=80,
94                        max_words=num_words_in_topic+1)
95         wc.generate_from_frequencies(words)
96
97         plt.clf()
98         plt.imshow(wc, interpolation="bilinear")
99         plt.axis("off")
100        plt.title(topic_name, y=-0.25, fontsize=20)
101        plt.gcf().tight_layout()
102        fig_path = 'img/{0}-{1}-{2}.png'.format(group, method, i+1, topic['topic'])
103        print('drawing ' + fig_path)
104        plt.savefig(fig_path)
105
106 def main():
107     cmd = sys.argv[1]
108     if cmd == 'show':
109         show_summary(load_pkl('blogs.pkl'))
110     elif cmd == 'eval':
111         fpath = sys.argv[2]
112         eval_topics(fpath, top_k=2)
113         eval_topics(fpath, top_k=2, method='tfidf')
114
115 if __name__ == '__main__':
116     main()

```